

# Features

Jong-Han Kim

EE787 Machine learning  
Kyung Hee University

# Records and embedding

## Raw data

- ▶ *raw data* pairs are  $(u, v)$ , with  $u \in \mathcal{U}$ ,  $v \in \mathcal{V}$
- ▶  $\mathcal{U}$  is set of all possible input values
- ▶  $\mathcal{V}$  is set of all possible output values
- ▶ each  $u$  is called a *record*
- ▶ typically a record is a tuple, or list,  $u = (u_1, u_2, \dots, u_r)$
- ▶ each  $u_i$  is a *field* or *component*, which has a *type*, e.g., real number, Boolean, categorical, ordinal, word, text, audio, image, parse tree (more on this later)
- ▶ e.g., a record for a house for sale might consist of  
(address, photo, description, house/apartment?, lot size,  $\dots$ , # bedrooms)

## Feature map

- ▶ learning algorithms are applied to  $(x, y)$  pairs,

$$x = \phi(u), \quad y = \psi(v)$$

- ▶  $\phi : \mathcal{U} \rightarrow \mathbf{R}^d$  is the *feature map* for  $u$
- ▶  $\psi : \mathcal{V} \rightarrow \mathbf{R}$  is the *feature map* for  $v$
- ▶ feature maps transform *records* into *vectors*
- ▶ feature maps usually work on each field separately,

$$\phi(u_1, \dots, u_r) = (\phi_1(u_1), \dots, \phi_r(u_r))$$


- ▶  $\phi_i$  is an *embedding* of the type of field  $i$  into a vector

# Embeddings

- ▶ embedding puts the different field types on an equal footing, *i.e.*, vectors
  - ▶ some embeddings are simple, *e.g.*,
    - ▶ for a number field ( $\mathcal{U} = \mathbf{R}$ ),  $\phi_i(u_i) = u_i$
    - ▶ for a Boolean field,  $\phi_i(u_i) = \begin{cases} 1 & u_i = \text{TRUE} \\ -1 & u_i = \text{FALSE} \end{cases}$
  - ▶ others are more sophisticated
    - ▶ text to TFID histogram
    - ▶ word2vec (maps words into vectors)
    - ▶ pre-trained ImageNet NN (maps images into vectors)
- (more on these later)

## More embeddings

- ▶ color to  $(R, G, B)$
- ▶ geolocation data:  $\phi(u) = (\text{Lat}, \text{Long})$  in  $\mathbf{R}^2$  or embed in  $\mathbf{R}^3$
- ▶ day of week:




## Faithful embeddings

a *faithful* embedding satisfies

- ▶  $\phi(u)$  is near  $\phi(\tilde{u})$  when  $u$  and  $\tilde{u}$  are 'similar'
- ▶  $\phi(u)$  is not near  $\phi(\tilde{u})$  when  $u$  and  $\tilde{u}$  are 'dissimilar'
- ▶ lefthand concept is *vector distance*
- ▶ righthand concept depends on field type, application
- ▶ interesting examples: names, professions, companies, countries, languages, ZIP codes, cities, songs, movies
- ▶ we will see later how such embeddings can be constructed

## Example: word2vec





## Standardized embeddings

usually assume that an embedding is *standardized*

- ▶ entries of  $\phi(u)$  are centered around 0
- ▶ entries of  $\phi(u)$  have RMS value around 1
- ▶ roughly speaking, entries of  $\phi(u)$  ranges over  $\pm 1$
- ▶ with standardized embeddings, entries of feature map

$$\phi(u_1, \dots, u_r) = (\phi_1(u_1), \dots, \phi_r(u_r))$$

are all comparable, *i.e.*, centered around zero, standard deviation around one

- ▶  $\text{rms}(\phi(u) - \phi(\tilde{u}))$  is reasonable measure of how close records  $u$  and  $\tilde{u}$  are

## Standardization or *z*-scoring

- ▶ suppose  $\mathcal{U} = \mathbf{R}$  (field type is real numbers)
- ▶ for data set  $u^1, \dots, u^n \in \mathbf{R}$

$$\bar{u} = \frac{1}{n} \sum_{i=1}^n u^i \quad \text{std}(u) = \left( \frac{1}{n} \sum_{i=1}^n (u^i - \bar{u})^2 \right)^{\frac{1}{2}}$$

- ▶ the *z-score* or *standardization* of  $u$  is the embedding

$$x = \text{zscore}(u) = \frac{1}{\text{std}(u)} (u - \bar{u})$$

- ▶ ensures that embedding values are centered at zero, with standard deviation one
- ▶ *z*-scored features are very easy to interpret:  $x = \phi(u) = +1.3$  means that  $u$  is 1.3 standard deviations above the mean value

## Standardized data matrix

- ▶ suppose all  $d$  (real) features have been standardized
- ▶ columns of  $n \times d$  feature matrix  $X$  have zero mean, RMS value one
- ▶  $(1/n)X^T X = \Sigma$  is the *feature correlation matrix*
- ▶  $\Sigma_{ii} = 1$  (since each column of  $X$  has RMS value 1, and so norm  $\sqrt{n}$ )
- ▶  $\Sigma_{ij}$  is *correlation coefficient* of  $i$ th and  $j$ th raw features

## Log transform

- ▶ old school rule-of-thumb: if field  $u$  is positive and ranges over wide scale, embed as  $\phi(u) = \log u$  (or  $\log(1 + u)$ ) (and then standarize)
- ▶ examples: web site visits, ad views, company capitalization
- ▶ interpretation as faithful embedding:
  - ▶ 20 and 22 are similar, as are 1000 and 1100
  - ▶ but 20 and 120 are not similar
  - ▶ *i.e.*, you care about fractional or relative differences between raw values(here, log embedding is faithful, affine embedding is not)
- ▶ can also apply to output or label field, *i.e.*,  $y = \psi(v) = \log v$  if you care about percentage or fractional errors; recover  $\hat{v} = \exp(\hat{y})$

## Example: House price prediction

- ▶ we want to predict house selling price  $v$  from record  $u = (u_1, u_2)$ 
  - ▶  $u_1 = \text{area (sq. ft.)}$
  - ▶  $u_2 = \# \text{ bedrooms}$
- ▶ we care about relative error in price, so we embed  $v$  as  $\psi(v) = \log v$  (and then standardize)
- ▶ we standardize fields  $u_1$  and  $u_2$

$$x_1 = \frac{u_1 - \mu_1}{\sigma_1}, \quad x_2 = \frac{u_2 - \mu_2}{\sigma_2}$$

- ▶  $\mu_1 = \bar{u}_1$  is mean area
- ▶  $\mu_2 = \bar{u}_2$  is mean number of bedrooms
- ▶  $\sigma_1 = \text{std}(u_1)$  is std. dev. of area
- ▶  $\sigma_2 = \text{std}(u_2)$  is std. dec. of # bedrooms

(means and std. dev. are over our data set)

## Example: House price regression model

► regression model:  $\hat{y} = \theta_1 + \theta_2 x_1 + \theta_3 x_2$

► in terms of original raw data:

$$\hat{v} = \exp \left( \theta_1 + \theta_2 \frac{u_1 - \mu_1}{\sigma_1} + \theta_3 \frac{u_2 - \mu_2}{\sigma_2} \right)$$

► exp undoes log embedding of house price

# Vector embeddings

## Vector embeddings for real field

- ▶ we can embed a field  $u$  into a vector  $x = \phi(u) \in \mathbf{R}^k$
- ▶ useful even when  $\mathcal{U} = \mathbf{R}$  (real field)
- ▶ polynomial embedding:

$$\phi(u) = (1, u, u^2, \dots, u^d)$$

- ▶ piecewise linear embedding:

$$\phi(u) = (1, (u)_-, (u)_+)$$

where  $(u)_- = \min(u, 0)$ ,  $(u)_+ = \max(u, 0)$


- ▶ regression with these features yield polynomial and piecewise linear predictors



## Whitening

- ▶ analog of standardization for raw data  $\mathcal{U} = \mathbf{R}^d$
- ▶ start with raw data,  $n \times d$  matrix  $U$
- ▶  $\bar{u} = U^T \mathbf{1}/n$  is vector of column means
- ▶  $\tilde{U} = U - \mathbf{1}\bar{u}^T$  is de-meanned data matrix
- ▶  $\tilde{U} = QR$  is its QR factorization
- ▶  $X = \sqrt{n}Q = \sqrt{n}\tilde{U}R^{-1}$  defines embedding  $x^i = \phi(u^i)$ 
  - ▶ columns of  $X$  have zero mean and RMS value one
  - ▶ columns of  $X$  are orthogonal
  - ▶ features are uncorrelated
  - ▶ feature correlation matrix is  $\Sigma = I$

## Whitening example



## Categorical data

- ▶ data field is *categorical* if it only takes a finite number of values
- ▶ i.e.,  $\mathcal{U}$  is a finite set  $\{\alpha_1, \dots, \alpha_k\}$
- ▶ examples:
  - ▶ TRUE/FALSE (two values, also called Boolean)
  - ▶ APPLE, ORANGE, BANANA (three values)
  - ▶ MONDAY, ..., SUNDAY (seven values)
  - ▶ ZIP code (40000 values)
- ▶ *one-hot embedding for categoricals*:  $\phi(\alpha_i) = e_i \in \mathbf{R}^k$   
$$\phi(\text{APPLE}) = (1, 0, 0), \quad \phi(\text{ORANGE}) = (0, 1, 0), \quad \phi(\text{BANANA}) = (0, 0, 1)$$
- ▶ standardizing these features handles *unbalanced* data

## Ordinal data

- ▶ ordinal data is categorical, with an order
- ▶ example: *Likert scale*, with values

STRONGLY DISAGREE, DISAGREE, NEUTRAL, AGREE, STRONGLY AGREE

- ▶ can embed into  $\mathbf{R}$  with values  $-2, -1, 0, 1, 2$
- ▶ or treat as categorical, with one-hot embedding into  $\mathbf{R}^5$
- ▶ example: number of bedrooms in house
  - ▶ can be treated as a real number
  - ▶ or as an ordinal with (say) values  $1, \dots, 6$

# Feature engineering

## How feature maps are constructed

- ▶ start by embedding each field

$$\phi(u_1, \dots, u_r) = (\phi_1(u_1), \dots, \phi_r(u_r))$$

- ▶ can then standardize, if needed
- ▶ use *feature engineering* to create new features from existing ones

## Creating new features

- ▶ product features:  $x_{\text{new}} = x_i x_j$  (models *interactions* between features)
- ▶ max features:  $x_{\text{new}} = \max(x_i, x_j)$  (can also use min)
- ▶ positive/negative parts:

$$x_{\text{new}+} = (x_i)_+ = \max(x_i, 0), \quad x_{\text{new}-} = (x_i)_- = \min(x_i, 0)$$

- ▶ random features:
  - ▶ choose random matrix  $R$
  - ▶ new features are  $(Rx)_+$  or  $(Rx)_-$

# Un-embedding



## Un-embedding

- ▶ we embed  $v$  as  $y = \psi(v)$ ,  $\psi : \mathcal{V} \rightarrow \mathbf{R}$
- ▶ we need to ‘invert’ this operation, and go from  $\hat{y}$  to  $\hat{v}$
- ▶ when the inverse function exists, we use  $\psi^{-1} : \mathbf{R} \rightarrow \mathcal{V}$
- ▶ example: log embedding  $y = \log v$  has inverse  $v = \exp y$
- ▶ prediction stack:
  1. *embed*: given record  $u$ , feature vector is  $x = \phi(u)$
  2. *predict*:  $\hat{y} = g(x)$
  3. *un-embed*:  $\hat{v} = \psi^{-1}(\hat{y})$
- ▶ final predictor is  $\hat{v} = \psi^{-1}(g(\phi(u)))$

## Un-embedding

- ▶ in many cases, the inverse of  $\psi$  function doesn't exist
- ▶ for example, embedding a Boolean or ordinal into  $\mathbf{R}$
- ▶ for the purposes of un-embedding, we define

$$\psi^{-1}(y) = \operatorname{argmin}_{v \in \mathcal{V}} \|y - \psi(v)\|$$

*i.e.*, we choose the value of  $v$  for which  $\psi(v)$  is closest to  $y$

- ▶ example: embed  $\text{TRUE} \mapsto 1$  and  $\text{FALSE} \mapsto -1$
- ▶ un-embed via

$$\psi^{-1}(y) = \begin{cases} \text{TRUE} & \text{if } y > 0 \\ \text{FALSE} & \text{otherwise} \end{cases}$$

## Example: Un-embedding one-hot

► *one-hot embedding*:  $\phi(u) = e_u$  for  $\mathcal{U} = \{1, \dots, d\}$

► un-embed

$$\phi^{-1}(x) = \underset{u}{\operatorname{argmin}} \|x - e_u\|_2 = \underset{u}{\operatorname{argmax}} x_u$$