

A short course on optimization for distributed machine learning

Jong-Han Kim

EIC7024
Kyung Hee University


Course contents

- ▶ Machine learning problems
 - ▶ Regression
 - ▶ Classification
- ▶ Optimization for machine learning
 - ▶ Optimization algorithms
 - ▶ Distributed optimization and statistical learning

Predictors

Information flow

raw data *feature* *outcome* *raw output*



student

hrs. sleep
hrs. study
hrs. facebook
alcohol freq.
⋮

exam score

grade

Data fitting

- ▶ we think $y \in \mathbf{R}$ and $x \in \mathbf{R}^d$ are (approximately) related by

$$y \approx f(x)$$

- ▶ x is called the *independent variable* or *feature vector*
- ▶ y is called the *outcome* or *response* or *target* or *label* or *dependent variable*
- ▶ often y is something we want to predict
- ▶ we don't know the 'true' relationship between x and y

Features

often x is a vector of features:

- ▶ documents
 - ▶ x is word count histogram for a document
- ▶ patient data
 - ▶ x are patient attributes, test results, symptoms
- ▶ customers
 - ▶ x is purchase history and other attributes of a customer

Where features come from

- ▶ we use u to denote the raw input data, such as a vector, word or text, image, video, audio, ...
- ▶ $x = \phi(u)$ is the corresponding *feature vector*
- ▶ the function ϕ is called the *embedding* or *feature function*
- ▶ ϕ might be very simple or quite complicated
- ▶ similarly, the raw output data v can be featurized as $y = \psi(v)$
- ▶ often we take $\phi(u)_1 = x_1 = 1$, the *constant feature*
- ▶ (much more on these ideas later)

Data and prior knowledge

- ▶ we are given data $x^1, \dots, x^n \in \mathbf{R}^d$ and $y^1, \dots, y^n \in \mathbf{R}$
- ▶ (x^i, y^i) is the i th *data pair* or *observation* or *example*
- ▶ we also (might) have *prior knowledge* about what f might look like
 - ▶ e.g., f is smooth or continuous: $f(x) \approx f(\tilde{x})$ when x is near \tilde{x}
 - ▶ or we might know $y \geq 0$

Predictor

- ▶ we seek a *predictor* or *model* $g : \mathbf{R}^d \rightarrow \mathbf{R}$
- ▶ for feature vector x , our prediction (of y) is $\hat{y} = g(x)$
- ▶ predictor g is chosen based on both data and prior knowledge
- ▶ in terms of raw data, our predictor is

$$\hat{v} = \psi^{-1}(g(\phi(u)))$$

(with a slight variation when ψ is not invertible)

- ▶ $\hat{y}^i \approx y^i$ means our predictor does well on i th data pair
- ▶ *but our real goal is to have $\hat{y} \approx y$ for (x, y) pairs we have not seen*

Prediction methods

- ▶ fraud, psychic powers, telepathy, magic sticks, incantations, crystals, hunches, statistics, AI, machine learning, data science
- ▶ and many algorithms . . .
- ▶ example: nearest neighbor predictor
 - ▶ given x , find its nearest neighbor x^i among given data
 - ▶ then predict $\hat{y} = g(x) = y^i$

A learning algorithm is a recipe for producing a predictor given data

Linear predictors

Linear predictor

- ▶ predictors that are linear functions of x are widely used
- ▶ a linear predictor has the form

$$g(x) = \theta^\top x$$

for some vector $\theta \in \mathbf{R}^d$, called the *predictor parameter vector*

- ▶ also called a *regression model*
- ▶ x_j is the j th feature, so the prediction is a linear combination of features

$$\hat{y} = g(x) = \theta_1 x_1 + \cdots + \theta_d x_d$$

- ▶ we get to choose the predictor parameter vector $\theta \in \mathbf{R}^d$
- ▶ sometimes we write $g_\theta(x)$ to emphasize the dependence on θ

Interpreting a linear predictor

$$\hat{y} = g(x) = \theta_1 x_1 + \cdots + \theta_d x_d$$

- ▶ θ_3 is the amount that prediction $\hat{y} = g(x)$ increases when x_3 increases by 1
 - ▶ particularly interpretable when x_3 is Boolean (only takes values 0 or 1)
- ▶ $\theta_7 = 0$ means that the prediction does not depend on x_7
- ▶ θ small means predictor is insensitive to changes in x :

$$|g(x) - g(\tilde{x})| = |\theta^\top x - \theta^\top \tilde{x}| = |\theta^\top (x - \tilde{x})| \leq \|\theta\| \|x - \tilde{x}\|$$

Affine predictor

- ▶ suppose the first feature is constant, $x_1 = 1$
- ▶ the linear predictor g is then an *affine function* of $x_{2:d}$, i.e., linear plus a constant

$$g(x) = \theta^\top x = \theta_1 + \theta_2 x_2 + \cdots + \theta_d x_d$$

- ▶ θ_1 is called the *offset* or *constant term* in the predictor
- ▶ θ_1 is the prediction when all features (except the constant) are zero

Empirical risk minimization

Loss function

a *loss* or *risk* function $\ell : \mathbf{R} \times \mathbf{R} \rightarrow \mathbf{R}$ quantifies how well (more accurately, how badly) \hat{y} approximates y

- ▶ smaller values of $\ell(\hat{y}, y)$ indicate that \hat{y} is a good approximation of y
- ▶ typically $\ell(y, y) = 0$ and $\ell(\hat{y}, y) \geq 0$ for all \hat{y}, y

examples

- ▶ *quadratic loss*: $\ell(\hat{y}, y) = (\hat{y} - y)^2$
- ▶ *absolute loss*: $\ell(\hat{y}, y) = |\hat{y} - y|$

Empirical risk

how well does the predictor g fit a data set (x^i, y^i) , $i = 1, \dots, n$, with loss ℓ ?

- ▶ the *empirical risk* is the average loss over the data points,

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n \ell(\hat{y}^i, y^i) = \frac{1}{n} \sum_{i=1}^n \ell(g(x^i), y^i)$$

- ▶ if \mathcal{L} is small, the predictor predicts the given data well
- ▶ when the predictor is parametrized by θ , we write

$$\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(g_{\theta}(x^i), y^i)$$

to show the dependence on the predictor parameter θ

Empirical risk minimization

- ▶ choosing the parameter θ in a parametrized predictor $g_\theta(x)$ is called *fitting the predictor* (to data)
- ▶ *empirical risk minimization (ERM)* is a general method for fitting a parametrized predictor
- ▶ ERM: *choose θ to minimize empirical risk $\mathcal{L}(\theta)$*
- ▶ thus, ERM chooses θ by attempting to match given data
- ▶ often there is no analytic solution to this minimization problem, so we use *numerical optimization* to find θ that minimizes $\mathcal{L}(\theta)$ (more on this topic later)

Least squares linear regression

Least squares linear regression

- ▶ linear predictor $\hat{y} = g_{\theta}(x) = \theta^{\top} x$
- ▶ $\theta \in \mathbf{R}^d$ is the model parameter
- ▶ we'll use square loss function $\ell(\hat{y}, y) = (\hat{y} - y)^2$
- ▶ empirical risk is MSE

$$\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n (\theta^{\top} x^i - y^i)^2$$

- ▶ ERM: choose model parameter θ to minimize MSE
- ▶ called *linear least squares fitting* or *linear regression*

Least squares formulation

- express MSE in matrix notation as

$$\begin{aligned}\mathcal{L}(\theta) &= \frac{1}{n} \sum_{i=1}^n (\theta^\top x^i - y^i)^2 = \frac{1}{n} \{ (\theta^\top x^1 - y^1)^2 + \dots + (\theta^\top x^n - y^n)^2 \} \\ &= \frac{1}{n} \left\| \begin{bmatrix} (x^1)^\top \theta - y^1 \\ \vdots \\ (x^n)^\top \theta - y^n \end{bmatrix} \right\|^2 \\ &= \frac{1}{n} \left\| \underbrace{\begin{bmatrix} (x^1)^\top \\ \vdots \\ (x^n)^\top \end{bmatrix}}_X \theta - \underbrace{\begin{bmatrix} y^1 \\ \vdots \\ y^n \end{bmatrix}}_y \right\|^2 = \frac{1}{n} \|X\theta - y\|^2\end{aligned}$$

- ERM is a *least squares problem*: choose θ to minimize $\|X\theta - y\|^2$
(factor $1/n$ doesn't affect choice of θ)

Least squares solution

- ▶ assuming X has linearly independent columns (which implies $n \geq d$), there is a unique optimal θ

$$\theta^* = (X^T X)^{-1} X^T y = X^\dagger y$$

- ▶ standard algorithm:
 - ▶ compute QR factorization $X = QR$ (e.g., Gram-Schmidt) (with orthogonal Q and invertible upper triangular R)
 - ▶ compute $Q^T y$
 - ▶ solve $R\theta^* = Q^T y$ by back substitution
- ▶ in Julia: `theta_opt = X \ y`
- ▶ complexity is $2d^2n$ flops

Data matrix

- ▶ the $n \times d$ matrix

$$X = \begin{bmatrix} (x^1)^\top \\ \vdots \\ (x^n)^\top \end{bmatrix}$$

is called the *data matrix*

- ▶ i th row of X is i th feature vector, transposed
- ▶ j th column of X gives values of j th feature x_j across our data set
- ▶ X_{ij} is the value of j th feature for the i th data point

Constant fit

- ▶ the simplest feature vector is constant: $x = \phi(u) = 1$
(doesn't depend on u !)
- ▶ corresponding predictor is a constant function: $g(x) = \theta_1$
- ▶ data matrix is $X = \mathbf{1}_n$
- ▶ so $X^\dagger = (X^\top X)^{-1} X^\top = (1/n) \mathbf{1}^\top$ and

$$\theta^* = X^\dagger y = \mathbf{1}^\top y / n = \text{avg}(y)$$

- ▶ *the average of the outcome values is the best constant predictor* (for square loss)
- ▶ optimal RMSE is standard deviation of outcome values

$$\left(\frac{1}{n} \sum_{i=1}^n (\text{avg}(y) - y^i)^2 \right)^{1/2}$$

Regression

- ▶ with $u \in \mathbf{R}^{d-1}$: $x = \phi(u) = (1, u)$
- ▶ same as $x_1 = 1$ (the first feature is constant)
- ▶ predictor has form

$$\hat{y} = \theta^\top x = \theta_1 + \theta_{2:d}^\top u$$

an affine function of u

Straight line fit

- ▶ with $u \in \mathbf{R}$, $x = (1, u) \in \mathbf{R}^2$
- ▶ model is $\hat{y} = g(x) = \theta_1 + \theta_2 u$
- ▶ this model is called *straight-line fit*
- ▶ when u is time, it's called the *trend line*
- ▶ when u is the whole market return, and y is an asset return, θ_2 is called ' β '

Constant versus straight-line fit models

- ▶ for the constant model, we choose θ_1 to minimize

$$\frac{1}{n} \sum_{i=1}^n (\theta_1 - y^i)^2$$

- ▶ for the straight-line model, we choose (θ_1, θ_2) to minimize

$$\frac{1}{n} \sum_{i=1}^n (\theta_1 + \theta_2 u^i - y^i)^2$$

- ▶ for optimal choices, this value is less than or equal to the one above (since we can take $\theta_2 = 0$ in the straight-line model)
- ▶ so the RMS error of the straight-line fit is no more than the standard deviation

Regularization

Sensitivity

- ▶ we have a linear predictor $\hat{y} = g(x) = \theta^\top x$
- ▶ if $|\theta_i|$ is large, then the prediction is very sensitive to x_i
(i.e., small changes in x_i lead to large changes in the prediction)
- ▶ large sensitivity can lead to overfit, poor generalization
(which would turn up in validation)
- ▶ for $x_1 = 1$ (the constant feature), there is no sensitivity, since the feature does not change
- ▶ suggests that we would like θ (or $\theta_{2:d}$ if $x_1 = 1$) not too large

Regularizer

- ▶ we will measure the size of θ using a *regularizer* function $r : \mathbf{R}^d \rightarrow \mathbf{R}$
- ▶ $r(\theta)$ is a measure of the size of θ (or $\theta_{2:d}$)

- ▶ *quadratic regularizer* (a.k.a. ℓ_2 or sum-of-squares):

$$r(\theta) = \|\theta\|^2 = \theta_1^2 + \dots + \theta_d^2$$

- ▶ *absolute value regularizer* (a.k.a. ℓ_1):

$$r(\theta) = \|\theta\|_1 = |\theta_1| + \dots + |\theta_d|$$

Regularized empirical risk minimization

- ▶ predictor should fit the given data well, *i.e.*, we want empirical risk

$$\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(\theta^\top x^i, y^i)$$

to be small

- ▶ predictor should not be too sensitive, *i.e.*, we want $r(\theta)$ small
- ▶ to trade off these two objectives, form *regularized empirical risk*

$$\mathcal{L}(\theta) + \lambda r(\theta)$$

where $\lambda \geq 0$ is the *regularization parameter* (or *hyper-parameter*)

- ▶ *regularized empirical risk minimization* (RERM): choose θ to minimize regularized empirical risk
- ▶ an optimization problem

Regularized empirical risk minimization

- ▶ for $\lambda = 0$, RERM reduces to ERM
- ▶ RERM produces a *family* of predictors, one for each value of λ
- ▶ in practice, we choose a few tens of values of λ , usually logarithmically spaced over a wide range
- ▶ use validation to choose among the candidate predictors
- ▶ we choose the largest value of λ that gives near minimum test error (*i.e.*, least sensitive predictor that generalizes well)

Ridge regression

- ▶ *ridge regression*: square loss and regularizer $r(\theta) = \|\theta\|^2$ (or $\|\theta_{2:d}\|^2$ if $x_1 = 1$)
- ▶ also called *Tykhonov regularized least squares*
- ▶ regularized empirical risk is

$$\begin{aligned}\mathcal{L}(\theta) + \lambda r(\theta) &= \|X\theta - y\|^2 + \lambda \|\theta\|^2 \\ &= \left\| \begin{bmatrix} X \\ \sqrt{\lambda}I \end{bmatrix} \theta - \begin{bmatrix} y \\ 0 \end{bmatrix} \right\|^2\end{aligned}$$

- ▶ so optimal θ is

$$\theta^* = \begin{bmatrix} X \\ \sqrt{\lambda}I \end{bmatrix}^\dagger \begin{bmatrix} y \\ 0 \end{bmatrix} = (X^T X + \lambda I)^{-1} X^T y$$

- ▶ (how do you modify this to handle $r(\theta) = \|\theta_{2:d}\|^2$?)

Penalty functions and error histograms

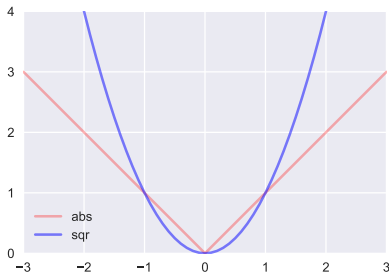
Loss and penalty functions

- ▶ empirical risk (or average loss) is $\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(\theta^T x^i, y^i)$
- ▶ the loss function $\ell(\hat{y}, y)$ penalizes deviation between the predicted value \hat{y} and the observed value y
- ▶ common form for loss function: $\ell(\hat{y}, y) = p(\hat{y} - y)$
- ▶ p is the *penalty function*
- ▶ e.g., the square penalty $p^{\text{sqr}}(r) = r^2$
- ▶ $r = \hat{y} - y$ is the *prediction error* or *residual*

Penalty functions

- ▶ the penalty function tells us how much we object to different values of prediction error
- ▶ usually $p(0) = 0$ and $p(r) \geq 0$ for all r
- ▶ if p is *symmetric*, i.e., $p(-r) = p(r)$, we care only about the magnitude (absolute value) of prediction error
- ▶ if p is *asymmetric*, i.e., $p(-r) \neq p(r)$, it bothers us more to over- or underestimate

Square versus absolute value penalty



- ▶ for square penalty $p^{\text{sqr}}(r) = r^2$
 - ▶ for small prediction errors, penalty is very small (small squared)
 - ▶ for large prediction errors, penalty is very large (large squared)
- ▶ for absolute penalty $p^{\text{abs}}(r) = |r|$
 - ▶ for small prediction errors, penalty is large (compared to square)
 - ▶ for large prediction errors, penalty is small (compared to square)

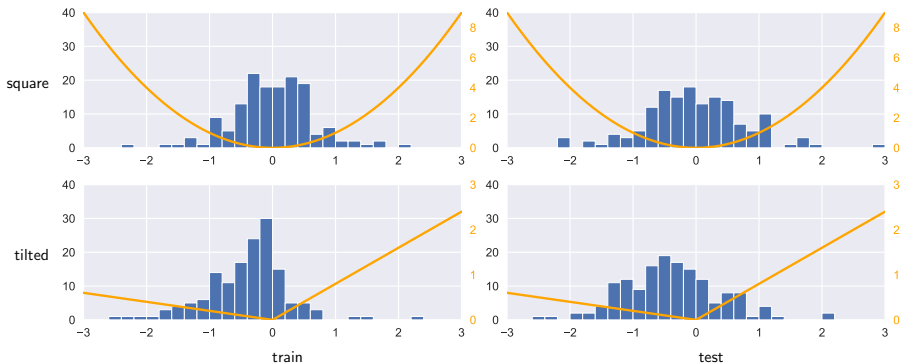
Predictors and choice of penalty function

- ▶ choice of penalty function depends on how you feel about large, small, positive, or negative prediction errors
- ▶ different choices of penalty function yield different predictor parameters
- ▶ choice of penalty function *shapes* the histogram of prediction errors, *i.e.*,

$$r^1, \dots, r^n$$

(usually divided into bins and displayed as bar graph distribution)

Histogram of residuals



- ▶ artificial data with $n = 300$ and $d = 30$, using 50/50 test/train split
- ▶ plots show histogram of residuals r^1, \dots, r^n
- ▶ tilted loss results in distribution with most residuals $r^i < 0$, i.e., predictor prefers $\hat{y}^i < y^i$

Robust fitting

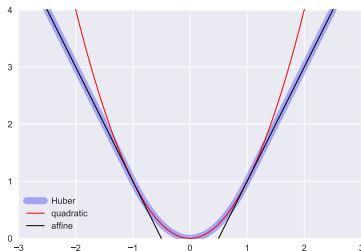
Outliers

- ▶ in some applications, a few data points are 'way off', or just 'wrong'
- ▶ occurs due to transcription errors, error in decimal point position, *etc.*
- ▶ these points are called *outliers*
- ▶ even a few outliers in a data set can result in a poor predictor
- ▶ several standard methods are used to remove outliers, or reduce their impact
- ▶ one simple method:
 - ▶ create predictor from data set
 - ▶ flag data points with large prediction errors as outliers
 - ▶ remove them from the data set and repeat

Robust penalty functions

- ▶ we say a penalty function is *robust* if it has low sensitivity to outliers
- ▶ robust penalty functions grow more slowly for large prediction error values than the square penalty
- ▶ and so 'allow' the predictor to have a few large prediction errors (presumably for the outliers)
- ▶ so they handle outliers more gracefully
- ▶ a *robust predictor* might fit, e.g., 98% of the data very well

Huber loss



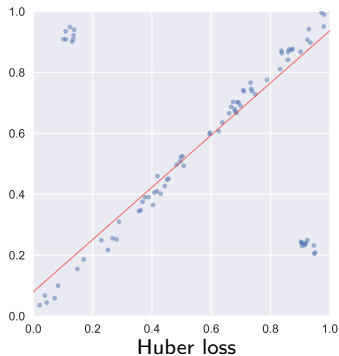
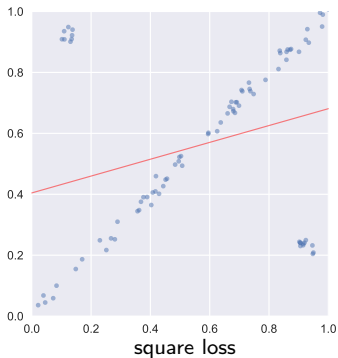
- ▶ the *Huber* penalty function is

$$p^{\text{hub}}(r) = \begin{cases} r^2 & \text{if } |r| \leq \alpha \\ \alpha(2|r| - \alpha) & \text{if } |r| > \alpha \end{cases}$$


- ▶ α is a parameter
- ▶ quadratic for small r , affine for large r

Huber loss

- ▶ linear growth for large r makes fit less sensitive to outliers
- ▶ ERM with Huber loss is called a *robust* prediction method



Log Huber

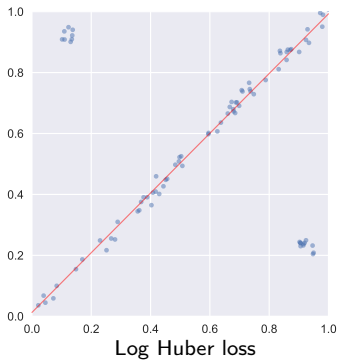


- quadratic for small y , logarithmic for large y

$$p^{\text{dh}}(y) = \begin{cases} y^2 & \text{if } |y| \leq \alpha \\ \alpha^2(1 - 2 \log(\alpha) + \log(y^2)) & \text{if } |y| > \alpha \end{cases}$$

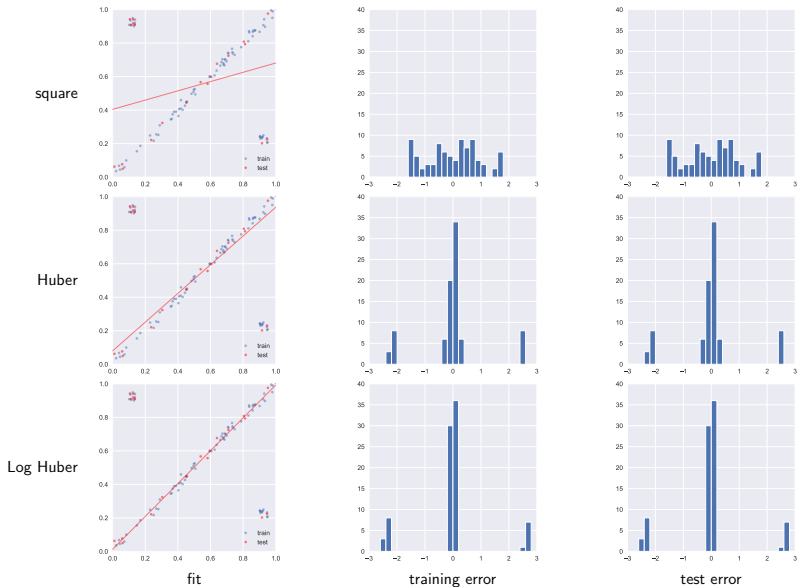
- diminishing incremental penalty at large y

Log Huber



- even less sensitive to outliers than Huber

Error distribution



Regularizers

Regularizers

- ▶ motivation:
 - ▶ large θ_i makes prediction $\theta^T x$ sensitive to value of x_i
 - ▶ so we want θ (or $\theta_{2:d}$ if $x_1 = 1$) small
- ▶ regularizer $r : \mathbf{R}^d \rightarrow \mathbf{R}$ measures the size of θ
- ▶ usually regularizer is *separable*,

$$r(\theta) = q(\theta_1) + \cdots + q(\theta_d)$$

where $q : \mathbf{R} \rightarrow \mathbf{R}$ is a penalty function for the predictor coefficients

Sum squares regularizer

- ▶ *sum squares* regularizer uses square penalty $q^{\text{sqr}}(a) = a^2$

$$r(\theta) = \|\theta\|^2 = \theta_1^2 + \cdots + \theta_d^2$$

- ▶ also called *quadratic*, *Tychonov*, or ℓ_2 regularizer

ℓ_1 regularizer

- ▶ *sum absolute* or ℓ_1 regularizer uses absolute value penalty $q^{\text{abs}}(a) = |a|$

$$r(\theta) = \|\theta\|_1 = |\theta_1| + \cdots + |\theta_d|$$

- ▶ $\|\theta\|_1$ is ℓ_1 *norm* of θ
- ▶ like the Euclidean or ℓ_2 norm $\|\theta\|$, it is a norm, *i.e.*, a measure of the size of the vector θ
- ▶ Euclidean norm is often written as $\|\theta\|_2$ to distinguish it from the ℓ_1 norm
- ▶ they are both members of the *p-norm family*, defined as

$$\|\theta\|_p = (|\theta_1|^p + \cdots + |\theta_d|^p)^{1/p}$$

for $p \geq 1$

Lasso regression

- ▶ use square loss $\ell(\hat{y}, y) = (\hat{y} - y)^2$
- ▶ choosing θ to minimize $\mathcal{L}(\theta) + \lambda \|\theta\|_2^2$ is called *ridge regression*
- ▶ choosing θ to minimize $\mathcal{L}(\theta) + \lambda \|\theta\|_1$ is called *lasso regression*
- ▶ invented by (Stanford's) Rob Tibshirani, 1994
- ▶ widely used in advanced machine learning
- ▶ unlike ridge regression, there is no formula for the lasso parameter vector
- ▶ but we can efficiently compute it anyway (since it's convex)
- ▶ the lasso regression model has some interesting properties

Sparsifying regularizers

Sparse coefficient vector

- ▶ suppose θ is sparse, *i.e.*, many of its entries are zero
- ▶ prediction $\theta^T x$ does not depend on features x_i for which $\theta_i = 0$
- ▶ this means we select *some* features to use (*i.e.*, those with $\theta_i \neq 0$)
- ▶ (possible) practical benefits of sparse θ :
 - ▶ can improve performance when many features are actually irrelevant
 - ▶ makes predictor *simpler to interpret*

Sparse coefficient vectors via ℓ_1 regularization

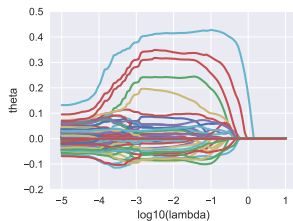
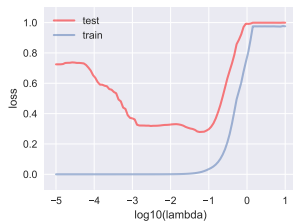
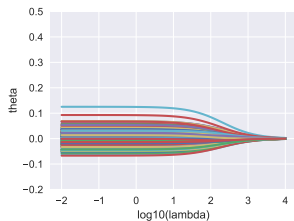
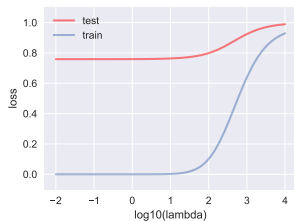
using ℓ_1 regularization leads to sparse coefficient vectors

$r(\theta) = \|\theta\|_1$ is called a *sparsifying regularizer*

rough explanation:

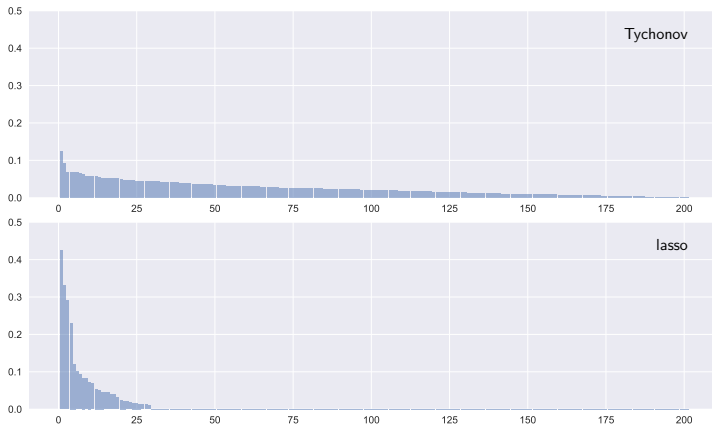
- ▶ for square penalty, once θ_i is small, θ_i^2 is very small
- ▶ so incentive for sum squares regularizer to make a coefficient smaller decreases once it is small
- ▶ for absolute penalty, incentive to make θ_i smaller keeps up all the way until it's zero

Example







- ▶ artificially generated 50 data points, 200 features
- ▶ only a few features are relevant
- ▶ left hand plots use Tychonov, right hand use lasso

Example



- ▶ sorted $|\theta_i|$ at optimal λ
- ▶ lasso solution has only 35 non-zero components

Example



- choose λ based on regularization path with test data
- keep features corresponding to largest components of θ and *retrain*
- plots above use most important 7 features identified by lasso

Even stronger sparsifiers

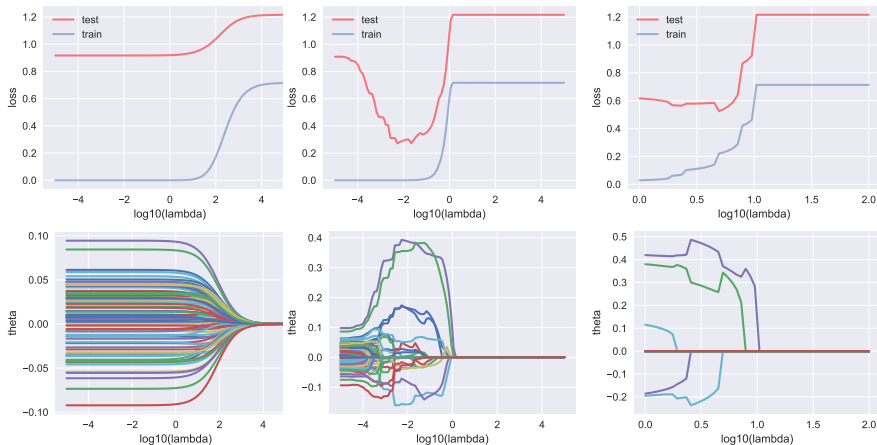
- ▶ $q(a) = |a|^{1/2}$
- ▶ called $\ell_{0.5}$ regularizer
- ▶ but you shouldn't use this term since

$$\left(|\theta_1|^{0.5} + \dots + |\theta_d|^{0.5}\right)^2$$

is not a norm (see VMLS)

- ▶ 'stronger' sparsifier than ℓ_1
- ▶ but not convex so computing θ is heuristic

Example



► ℓ_2 , ℓ_1 , and square root regularization

Nonnegative regularizer

Nonnegative coefficients

- ▶ in some cases we know or require that $\theta_i \geq 0$
- ▶ this means that when x_i increases, so must our prediction
- ▶ we can think of this constraint as regularization with penalty function

$$q(a) = \begin{cases} 0 & a \geq 0 \\ \infty & a < 0 \end{cases}$$

- ▶ example: y is lifespan, x_i measures healthy behavior i
- ▶ with quadratic loss, called *nonnegative least squares* (NNLS)
- ▶ common heuristic for nonnegative least squares: use $(\theta^{\text{ls}})_+$ (works poorly)

Classification

Boolean classification

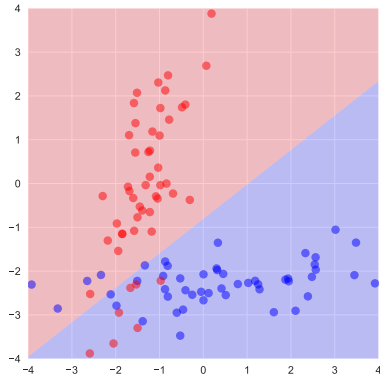
- ▶ supervised learning is called *boolean classification* when raw output variable v is a categorical that can take two possible values
- ▶ we denote these -1 and 1 , and they often correspond to $\{\text{FALSE}, \text{TRUE}\}$ or $\{\text{NEGATIVE}, \text{POSITIVE}\}$
- ▶ for a data record u^i, v^i , the value $v^i \in \{-1, 1\}$ is called the *class* or *label*
- ▶ a *boolean classifier* predicts label \hat{v} given raw input u

Classification



- ▶ here $u \in \mathbb{R}^2$
- ▶ red points have $v^i = -1$, blue points have $v^i = 1$
- ▶ we'd like a predictor that maps any $u \in \mathbb{R}^2$ into prediction $\hat{v} \in \{-1, 1\}$

Example: Least squares classifier



- ▶ embed $x = (1, u)$ and $y = v$, apply least squares regression
- ▶ gives $\hat{y} = \theta_1 + \theta_2 u_1 + \theta_3 u_2$
- ▶ predict using $\hat{v} = \text{sign}(\hat{y})$
- ▶ 11% of points misclassified at training

Confusion matrix

The two types of errors

- ▶ measure performance of a specific predictor on a set of n data records
- ▶ each data point i has $v^i \in \{-1, 1\}$
- ▶ and corresponding prediction $\hat{v}^i = g(v^i) \in \{-1, 1\}$
- ▶ only four possible values for the data pair \hat{v}^i, v^i :
 - ▶ *true positive* if $\hat{v}^i = 1$ and $v^i = 1$
 - ▶ *true negative* if $\hat{v}^i = -1$ and $v^i = -1$
 - ▶ *false negative* or *type II error* if $\hat{v}^i = -1$ and $v^i = 1$
 - ▶ *false positive* or *type I error* if $\hat{v}^i = 1$ and $v^i = -1$

Confusion matrix

- ▶ for a predictor and a data set define the *confusion matrix*

$$C = \begin{bmatrix} \# \text{ true negatives} & \# \text{ false negatives} \\ \# \text{ false positives} & \# \text{ true positives} \end{bmatrix} = \begin{bmatrix} C_{\text{tn}} & C_{\text{fn}} \\ C_{\text{fp}} & C_{\text{tp}} \end{bmatrix}$$

(warning: some people use the transpose of C)

- ▶ $C_{\text{tn}} + C_{\text{fn}} + C_{\text{fp}} + C_{\text{tp}} = n$ (total number of examples)
- ▶ $N_{\text{n}} = C_{\text{tn}} + C_{\text{fp}}$ is number of negative examples
- ▶ $N_{\text{p}} = C_{\text{fn}} + C_{\text{tp}}$ is number of positive examples
- ▶ diagonal entries give numbers of correct predictions
- ▶ off-diagonal entries give numbers of incorrect predictions of the two types

Some boolean classification measures

- ▶ confusion matrix $\begin{bmatrix} C_{tn} & C_{fn} \\ C_{fp} & C_{tp} \end{bmatrix}$
- ▶ the basic error measures:
 - ▶ *false positive rate* is C_{fp}/n
 - ▶ *false negative rate* is C_{fn}/n
 - ▶ *error rate* is $(C_{fn} + C_{fp})/n$
- ▶ error measures some people use:
 - ▶ *true positive rate* or *sensitivity* or *recall* is C_{tp}/N_p
 - ▶ *false alarm rate* is C_{fp}/N_n
 - ▶ *specificity* or *true negative rate* is C_{tn}/N_n
 - ▶ *precision* is $C_{tp}/(C_{tp} + C_{fp})$

Neyman-Pearson error

- ▶ *Neyman-Pearson error* over a data set is $\kappa C_{\text{fn}}/n + C_{\text{fp}}/n$
 - ▶ a scalarization of our two objectives, false positive and false negative rates
 - ▶ κ is how much more false negatives irritate us than false positives
 - ▶ when $\kappa = 1$, the Neyman-Pearson error is the *error rate*
-
- ▶ we'll use the Neyman-Pearson error as our scalarized measure

ERM for classification tasks

Embedding

- ▶ we embed raw input and output records as $x = \phi(u)$ and $y = \psi(v)$
- ▶ ϕ is the feature map
- ▶ ψ is *the identity map*, $\psi(v) = v$
- ▶ un-embed by $\hat{v} = \text{sign}(\hat{y})$
- ▶ equivalent to $\hat{v} = \underset{v \in \{-1, 1\}}{\text{argmin}} |\hat{y} - \psi(v)|$
- ▶ i.e., choose the nearest boolean value to the (real) prediction \hat{y}

ERM

- ▶ given loss function $\ell(\hat{y}, y)$, *empirical risk* on a data set is

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n \ell(\hat{y}^i, y^i)$$

- ▶ for linear model $\hat{y} = \theta^\top x$, with $\theta \in \mathbf{R}^d$,

$$\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(\theta^\top x^i, y^i)$$

- ▶ ERM: choose θ to minimize $\mathcal{L}(\theta)$
- ▶ regularized ERM: choose θ to minimize $\mathcal{L}(\theta) + \lambda r(\theta)$, with $\lambda > 0$

Loss functions for boolean classification

- ▶ to apply ERM, we need a loss function on embedded variables $\ell(\hat{y}, y)$
- ▶ y can only take values -1 or 1
- ▶ but $\hat{y} = \theta^\top x \in \mathbf{R}$ can be any real number
- ▶ to specify ℓ , we only need to give two functions (of a scalar \hat{y}):
 - ▶ $\ell(\hat{y}, -1)$ is how much \hat{y} irritates us when $y = -1$
 - ▶ $\ell(\hat{y}, 1)$ is how much \hat{y} irritates us when $y = 1$
- ▶ we can take $\ell(\hat{y}, 1) = \kappa \ell(-\hat{y}, -1)$, to reflect that false negatives irritate us a factor κ more than false positives

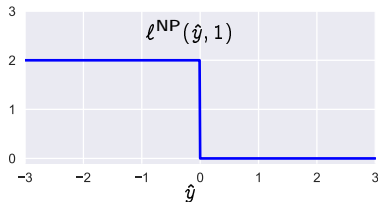
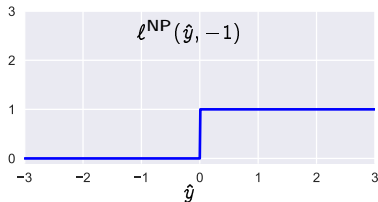
Neyman-Pearson loss

► Neyman-Pearson loss is

$$\ell^{\text{NP}}(\hat{y}, -1) = \begin{cases} 1 & \hat{y} \geq 0 \\ 0 & \hat{y} < 0 \end{cases}$$

$$\ell^{\text{NP}}(\hat{y}, 1) = \kappa \ell^{\text{NP}}(\hat{y}, -1) = \begin{cases} \kappa & \hat{y} < 0 \\ 0 & \hat{y} \geq 0 \end{cases}$$

► empirical Neyman-Pearson risk \mathcal{L}^{NP} is the Neyman-Pearson error



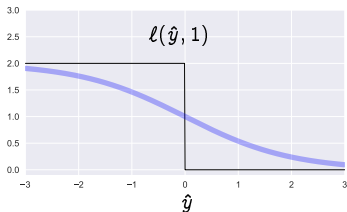
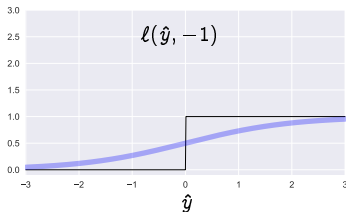
The problem with Neyman-Pearson loss

- ▶ empirical Neyman-Pearson risk $\mathcal{L}^{\text{NP}}(\theta)$ is not differentiable, or even continuous (and certainly not convex)
- ▶ worse, its gradient $\nabla \mathcal{L}^{\text{NP}}(\theta)$ is either zero or undefined
- ▶ so an optimizer does not know how to improve the predictor

Idea of proxy loss


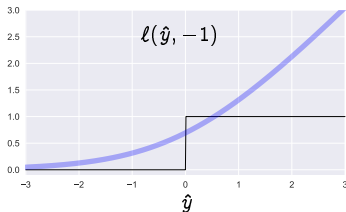
- ▶ we get better results using a *proxy loss* that
 - ▶ approximates, or at least captures the flavor of, the Neyman-Pearson loss
 - ▶ is more easily optimized (e.g., is convex or has nonzero derivative)
- ▶ we want a proxy loss function
 - ▶ with $\ell(\hat{y}, -1)$ small when $\hat{y} < 0$, and larger when $\hat{y} > 0$
 - ▶ with $\ell(\hat{y}, +1)$ small when $\hat{y} > 0$, and larger when $\hat{y} < 0$
 - ▶ which has other nice characteristics, e.g., differentiable or convex

Sigmoid loss



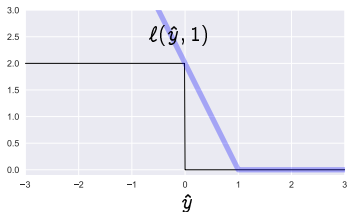
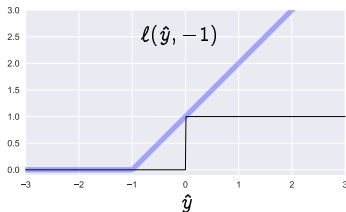
- ▶ $\ell(\hat{y}, -1) = \frac{1}{1 + e^{-\hat{y}}}$, $\ell(\hat{y}, 1) = \kappa \ell(-\hat{y}, -1) = \frac{\kappa}{1 + e^{\hat{y}}}$
- ▶ differentiable approximation of Neyman-Pearson loss
- ▶ but not convex

Logistic loss



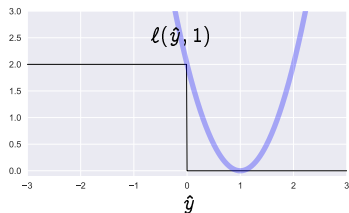
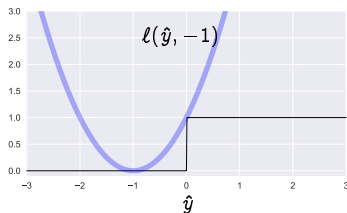
- ▶ $\ell(\hat{y}, -1) = \log(1 + e^{\hat{y}})$, $\ell(\hat{y}, 1) = \kappa \ell(-\hat{y}, -1) = \kappa \log(1 + e^{-\hat{y}})$
- ▶ differentiable and convex approximation of Neyman-Pearson loss

Hinge loss



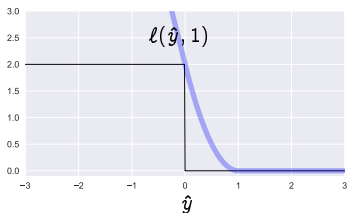
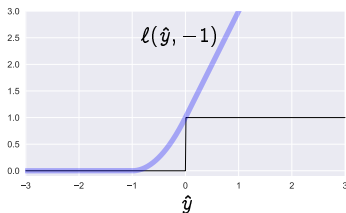
- ▶ $\ell(\hat{y}, -1) = (1 + \hat{y})_+$, $\ell(\hat{y}, 1) = \kappa \ell(-\hat{y}, -1) = \kappa(1 - \hat{y})_+$
- ▶ nondifferentiable but convex approximation of Neyman-Pearson loss

Square loss



- ▶ $\ell(\hat{y}, -1) = (1 + \hat{y})^2$, $\ell(\hat{y}, 1) = \kappa \ell(-\hat{y}, -1) = \kappa(1 - \hat{y})^2$
- ▶ ERM is least squares problem

Hubristic loss



- define the *hubristic loss* (huber + logistic) as

$$\ell(\hat{y}, -1) = \begin{cases} 0 & \hat{y} < -1 \\ (\hat{y} + 1)^2 & -1 \leq \hat{y} \leq 0 \\ 1 + 2\hat{y} & \hat{y} > 0 \end{cases}$$

- $\ell(\hat{y}, 1) = \kappa \ell(-\hat{y}, -1)$

Boolean classifiers

Least squares classifier

- ▶ use empirical risk with square loss

$$\mathcal{L}(\theta) = \frac{1}{n} \left(\sum_{i: y^i = -1} (1 + \hat{y}^i)^2 + \kappa \sum_{i: y^i = 1} (1 - \hat{y}^i)^2 \right)$$

and your choice of regularizer

- ▶ with sum squares regularizer, this is *least squares classifier*
- ▶ we can minimize $\mathcal{L}(\theta) + \lambda r(\theta)$ using, e.g., QR factorization

Logistic regression

- ▶ use empirical risk with logistic loss

$$\mathcal{L}(\theta) = \frac{1}{n} \left(\sum_{i: y^i = -1} \log(1 + e^{\hat{y}^i}) + \kappa \sum_{i: y^i = 1} \log(1 + e^{-\hat{y}^i}) \right)$$

and your choice of regularizer

- ▶ can minimize $\mathcal{L}(\theta) + \lambda r(\theta)$ using prox-gradient method
- ▶ we will find an actual minimizer if r is convex

Support vector machine

(usually abbreviated as *SVM*)

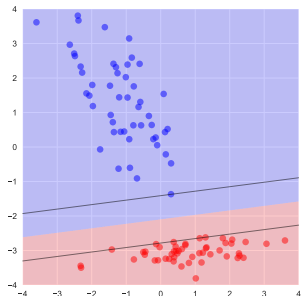
- use empirical risk with hinge loss

$$\mathcal{L}(\theta) = \frac{1}{n} \left(\sum_{i: y^i = -1} (1 + \hat{y}^i)_+ + \kappa \sum_{i: y^i = 1} (1 - \hat{y}^i)_+ \right)$$

and sum squares regularizer

- $\mathcal{L}(\theta) + \lambda r(\theta)$ is convex
- it can be minimized by various methods (but not prox-gradient)

Support vector machine



- decision boundary is $\theta^T x = 0$
- black lines show points where $\theta^T x = \pm 1$
- what is the training risk here?

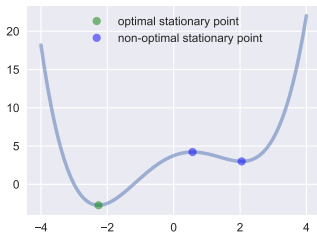
Optimization problems and algorithms

Optimization problem

$$\text{minimize } f(\theta)$$

- ▶ $\theta \in \mathbf{R}^d$ is the *variable* or *decision variable*
- ▶ $f : \mathbf{R}^d \rightarrow \mathbf{R}$ is the *objective function*
- ▶ goal is to choose θ to minimize f
- ▶ θ^* is *optimal* means that for all θ , $f(\theta) \geq f(\theta^*)$
- ▶ $f^* = f(\theta^*)$ is the *optimal value* of the problem
- ▶ optimization problems arise in many fields and applications, including machine learning

Optimality condition



- ▶ let's assume that f is *differentiable*, i.e., partial derivatives $\frac{\partial f(\theta)}{\partial \theta_i}$ exist
- ▶ if θ^* is optimal, then $\nabla f(\theta^*) = 0$
- ▶ $\nabla f(\theta) = 0$ is called the *optimality condition* for the problem
- ▶ there can be points that satisfy $\nabla f(\theta) = 0$ but are not optimal
- ▶ we call points that satisfy $\nabla f(\theta) = 0$ *stationary points*
- ▶ not all stationary points are optimal

Solving optimization problems

- ▶ in some cases, we can solve the problem analytically
- ▶ e.g., least squares: minimize $f(\theta) = \|X\theta - y\|^2$
 - ▶ optimality condition is $\nabla f(\theta) = 2X^T(X\theta - y) = 0$
 - ▶ this has (unique) solution $\theta^* = (X^T X)^{-1} X^T y = X^\dagger y$
(when columns of X are linearly independent)
- ▶ in other cases, we resort to an *iterative algorithm* that computes a sequence $\theta^1, \theta^2, \dots$ with, hopefully, $f(\theta^k) \rightarrow f^*$ as $k \rightarrow \infty$

Iterative algorithms

- ▶ *iterative algorithm* computes a sequence $\theta^1, \theta^2, \dots$
- ▶ θ^k is called the k th *iterate*
- ▶ θ^1 is called the *starting point*
- ▶ many iterative algorithms are *descent methods*, which means

$$f(\theta^{k+1}) < f(\theta^k), \quad k = 1, 2, \dots$$

i.e., each iterate is better than the previous one

- ▶ this means that $f(\theta^k)$ converges, but not necessarily to f^*



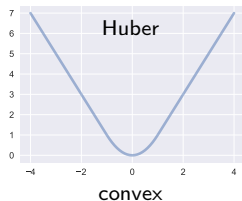
Stopping criterion

- ▶ in practice, we stop after a finite number K of steps
- ▶ typical stopping criterion: stop if $\|\nabla f(\theta^k)\| \leq \epsilon$ or $k = k^{\max}$
- ▶ ϵ is a small positive number, the *stopping tolerance*
- ▶ k^{\max} is the maximum number of iterations
- ▶ in words: we stop when θ^k is almost a stationary point
- ▶ we hope that $f(\theta^K)$ is not too much bigger than f^*
- ▶ or more realistically, that θ^K is at least useful for our application

Non-heuristic and heuristic algorithms

- ▶ in some cases we *know* that $f(\theta^k) \rightarrow f^*$, for any θ^1
 - ▶ in words: *we'll get to a solution if we keep iterating*
 - ▶ called *non-heuristic*
-
- ▶ other algorithms do not guarantee that $f(\theta^k) \rightarrow f^*$
 - ▶ we can hope that even if $f(\theta^k) \not\rightarrow f^*$, θ^k is still useful for our application
 - ▶ called *heuristic*

Convex functions



- ▶ a function $f : \mathbf{R}^d \rightarrow \mathbf{R}$ is **convex** if for any θ , $\tilde{\theta}$, and α with $0 \leq \alpha \leq 1$,

$$f(\alpha\theta + (1 - \alpha)\tilde{\theta}) \leq \alpha f(\theta) + (1 - \alpha)f(\tilde{\theta})$$

- ▶ roughly speaking, f has 'upward curvature'
- ▶ for $d = 1$, same as $f''(\theta) \geq 0$ for all θ

Convex optimization

- optimization problem

$$\text{minimize } f(\theta)$$

is called *convex* if the objective function f is convex

- for convex optimization problem, $\nabla f(\theta) = 0$ only for θ optimal, i.e.,
all stationary points are optimal

- algorithms for convex optimization are non-heuristic
- i.e., *we can solve convex optimization problems* (exactly, in principle)

Convex ERM problems

- regularized empirical risk function $f(\theta) = \mathcal{L}(\theta) + \lambda r(\theta)$, with $\lambda \geq 0$,

$$\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n p(\theta^\top x^i - y^i), \quad r(\theta) = q(\theta_1) + \cdots + q(\theta_d)$$

- f is convex if loss penalty p and parameter penalty q functions are convex
- convex penalties: square, absolute, tilted absolute, Huber
- non-convex penalties: log Huber, squareroot

Gradient method

Gradient method

- ▶ assume f is differentiable
- ▶ at iteration θ^k , create affine (Taylor) approximation of f valid near θ^k

$$\hat{f}(\theta; \theta^k) = f(\theta^k) + \nabla f(\theta^k)^T (\theta - \theta^k)$$

- ▶ $\hat{f}(\theta; \theta^k) \approx f(\theta)$ for θ near θ^k
- ▶ choose θ^{k+1} to make $\hat{f}(\theta^{k+1}; \theta^k)$ small, but with $\|\theta^{k+1} - \theta^k\|$ not too large
- ▶ choose θ^{k+1} to minimize $\hat{f}(\theta; \theta^k) + \frac{1}{2h^k} \|\theta - \theta^k\|^2$
- ▶ $h^k > 0$ is a *trust parameter* or *step length* or *learning rate*
- ▶ solution is $\theta^{k+1} = \theta^k - h^k \nabla f(\theta^k)$
- ▶ roughly: take step in direction of negative gradient

Gradient method update

- choose θ^{k+1} to as minimizer of

$$f(\theta^k) + \nabla f(\theta^k)^T(\theta - \theta^k) + \frac{1}{2h^k} \|\theta - \theta^k\|^2$$

- rewrite as

$$f(\theta^k) + \frac{1}{2h^k} \|(\theta - \theta^k) + h^k \nabla f(\theta^k)\|^2 - \frac{h^k}{2} \|\nabla f(\theta^k)\|^2$$

- first and third terms don't depend on θ
- middle term is minimized (made zero!) by choice

$$\theta = \theta^k - h^k \nabla f(\theta^k)$$

How to choose step length

- ▶ if h^k is too large, we can have $f(\theta^{k+1}) > f(\theta^k)$
- ▶ if h^k is too small, we have $f(\theta^{k+1}) < f(\theta^k)$ but progress is slow
- ▶ a simple scheme:
 - ▶ if $f(\theta^{k+1}) > f(\theta^k)$, set $h^{k+1} = h^k/2$, $\theta^{k+1} = \theta^k$ (a *rejected step*)
 - ▶ if $f(\theta^{k+1}) \leq f(\theta^k)$, set $h^{k+1} = 1.2h^k$ (an *accepted step*)
- ▶ reduce step length by half if it's too long; increase it 20% otherwise

Gradient method summary

choose an initial $\theta^1 \in \mathbf{R}^d$ and $h^1 > 0$ (e.g., $\theta^1 = 0$, $h^1 = 1$)

for $k = 1, 2, \dots, k^{\max}$

1. compute $\nabla f(\theta^k)$; quit if $\|\nabla f(\theta^k)\|$ is small enough
2. form tentative update $\theta^{\text{tent}} = \theta^k - h^k \nabla f(\theta^k)$
3. if $f(\theta^{\text{tent}}) \leq f(\theta^k)$, set $\theta^{k+1} = \theta^{\text{tent}}$, $h^{k+1} = 1.2h^k$
4. else set $h^k := 0.5h^k$ and go to step 2

Gradient method convergence

- ▶ (assuming some technical conditions hold) we have

$$\|\nabla f(\theta^k)\| \rightarrow 0 \text{ as } k \rightarrow \infty$$

- ▶ i.e., the gradient method always finds a stationary point

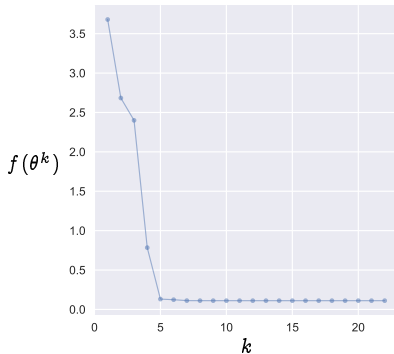
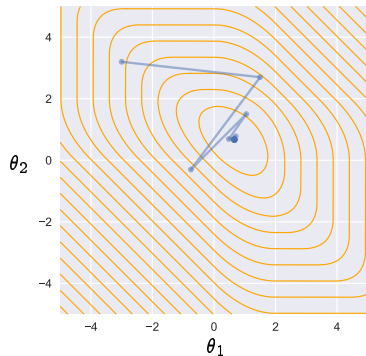
- ▶ for *convex problems*

- ▶ gradient method is *non-heuristic*
- ▶ for any starting point θ^1 , $f(\theta^k) \rightarrow f^*$ as $k \rightarrow \infty$

- ▶ for *non-convex problems*


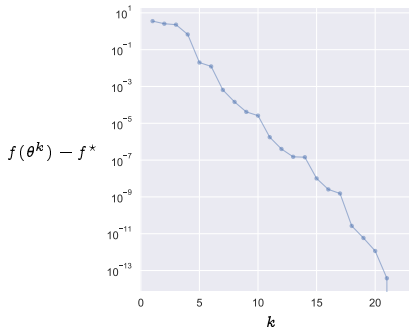
- ▶ gradient method is *heuristic*
- ▶ we can (and often do) have $f(\theta^k) \not\rightarrow f^*$

Example: Convex objective



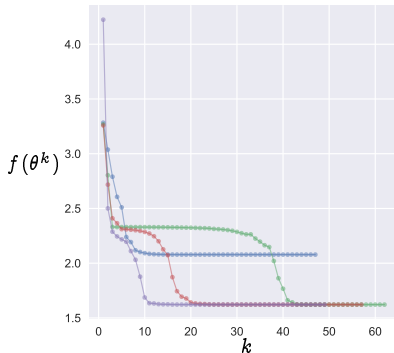
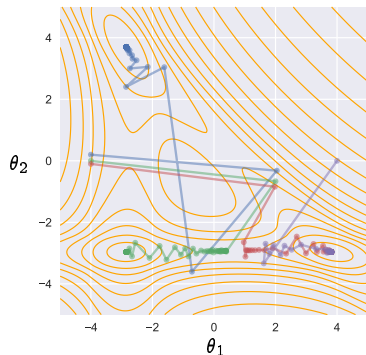
- ▶ $f(\theta) = \frac{1}{3} (p^{\text{hub}}(\theta_1 - 1) + p^{\text{hub}}(\theta_2 - 1) + p^{\text{hub}}(\theta_1 + \theta_2 - 1))$
- ▶ f is convex
- ▶ optimal point is $\theta^* = (2/3, 2/3)$, with $f^* = 1/9$

Example: Convex objective



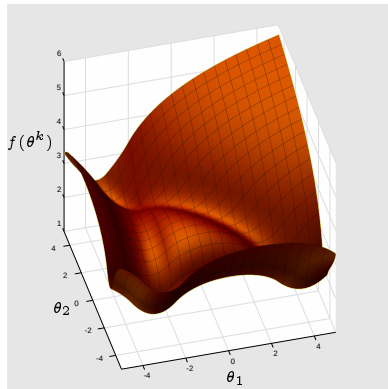
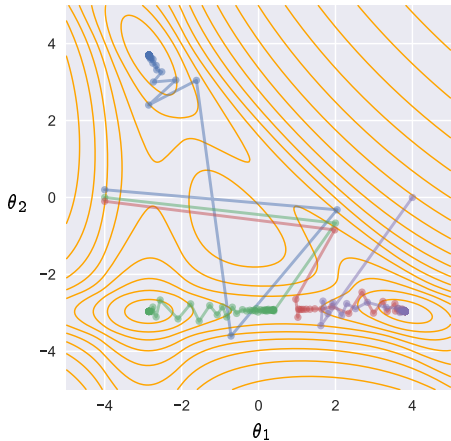
- $f(\theta^k)$ is a decreasing function of k , (roughly) exponentially
- $\|\nabla f(\theta^k)\| \rightarrow 0$ as $k \rightarrow \infty$

Example: Non-convex objective

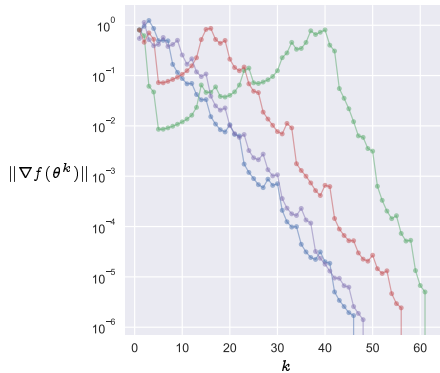
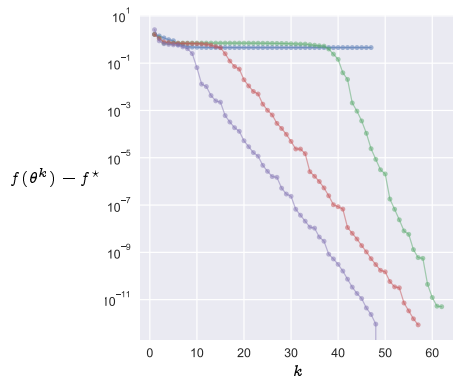


- ▶ $f(\theta) = \frac{1}{3} (p^{\text{lh}}(\theta_1 + 3) + p^{\text{lh}}(2\theta_2 + 6) + p^{\text{lh}}(\theta_1 + \theta_2 - 1))$
- ▶ f is sum of log-Huber functions, so not convex
- ▶ gradient algorithm converges, but limit depends on initial guess

Example: Non-convex objective



Example: Non-convex objective



Gradient method for ERM

Gradient of empirical risk function

- ▶ empirical risk is sum of terms for each data point

$$\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(\hat{y}^i, y^i) = \frac{1}{n} \sum_{i=1}^n \ell(\theta^T x^i, y^i)$$

- ▶ convex if loss function ℓ is convex in first argument
- ▶ gradient is sum of terms for each data point

$$\nabla \mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n \ell'(\theta^T x^i, y^i) x^i$$

where $\ell'(\hat{y}, y)$ is derivative of ℓ with respect to its first argument \hat{y}

Evaluating gradient of empirical risk function

- ▶ compute n -vector $\hat{y}^k = X\theta^k$
 - ▶ compute n -vector z^k , with entries $z_i^k = \ell'(\hat{y}_i^k, y^i)$
 - ▶ compute d -vector $\nabla \mathcal{L}(\theta^k) = (1/n)X^T z^k$
-
- ▶ first and third steps are matrix-vector multiplication, each costing $2nd$ flops
 - ▶ second step costs order n flops (dominated by other two)
 - ▶ total is $4nd$ flops

Prox-gradient method

Minimizing composite functions

- ▶ want to minimize $F(\theta) = f(\theta) + g(\theta)$ (called *composite function*)
- ▶ f is differentiable, but g need not be
- ▶ example: minimize $\mathcal{L}(\theta) + \lambda r(\theta)$, with $r(\theta) = \|\theta\|_1$
- ▶ we'll see idea of gradient method extends directly to composite functions

Selective linearization

- ▶ at iteration k , linearize f *but not* g

$$\hat{F}(\theta; \theta^k) = f(\theta^k) + \nabla f(\theta^k)^T(\theta - \theta^k) + g(\theta)$$

- ▶ want $\hat{F}(\theta; \theta^k)$ small, but with θ near θ^k
- ▶ choose θ^{k+1} to minimize $\hat{F}(\theta; \theta^k) + \frac{1}{2h^k}\|\theta - \theta^k\|^2$, with $h^k > 0$
- ▶ same as minimizing

$$g(\theta) + \frac{1}{2h^k}\|\theta - (\theta^k - h^k \nabla f(\theta^k))\|^2$$

- ▶ for many 'simple' functions g , this minimization can be done analytically
- ▶ this iteration from θ^k to θ^{k+1} is called *prox-gradient step*

Prox-gradient iteration

- prox-gradient iteration has two parts:

1. *gradient step*: $\theta^{k+1/2} = \theta^k - h^k \nabla f(\theta^k)$

2. *prox step*: choose θ^{k+1} to minimize $g(\theta) + \frac{1}{2h^k} \|\theta - \theta^{k+1/2}\|^2$

($\theta^{k+1/2}$ is an intermediate iterate, in between θ^k and θ^{k+1})

- step 1 handles differentiable part of objective, *i.e.*, f
- step 2 handles second part of objective, *i.e.*, g

Proximal operator

- ▶ given function $q : \mathbf{R}^d \rightarrow \mathbf{R}$, and $\kappa > 0$,

$$\mathbf{prox}_{q,\kappa}(v) = \underset{\theta}{\operatorname{argmin}} \left(q(\theta) + \frac{1}{2\kappa} \|\theta - v\|^2 \right)$$

is called the *proximal operator* of q at v , with parameter κ

- ▶ the prox-gradient step can be expressed as

$$\theta^{k+1} = \mathbf{prox}_{g,h^k}(\theta^{k+1/2}) = \mathbf{prox}_{g,h^k}(\theta^k - h^k \nabla f(\theta^k))$$

- ▶ hence the name prox-gradient iteration

How to choose step length

- ▶ same as for gradient, but using $F(\theta) = f(\theta) + g(\theta)$
- ▶ a simple scheme:
 - ▶ if $F(\theta^{k+1}) > F(\theta^k)$, set $h^{k+1} = h^k/2$, $\theta^{k+1} = \theta^k$ (a *rejected step*)
 - ▶ if $F(\theta^{k+1}) \leq F(\theta^k)$, set $h^{k+1} = 1.2h^k$ (an *accepted step*)
- ▶ reduce step length by half if it's too long; increase it 20% otherwise

Stopping criterion

- ▶ stopping condition for prox-gradient method:

$$\left\| \nabla f(\theta^{k+1}) - \frac{1}{h^k}(\theta^{k+1} - \theta^{k+1/2}) \right\| \leq \epsilon$$

- ▶ analog of $\|\nabla f(\theta^{k+1})\| \leq \epsilon$ for gradient method
- ▶ second term $-\frac{1}{h^k}(\theta^{k+1} - \theta^{k+1/2})$ serves the purpose of a gradient for g (which need not be differentiable)

Prox-gradient method summary

choose an initial $\theta^1 \in \mathbf{R}^d$ and $h^1 > 0$ (e.g., $\theta^1 = 0$, $h^1 = 1$)

for $k = 1, 2, \dots, k^{\max}$

1. gradient step. $\theta^{k+1/2} = \theta^k - h^k \nabla f(\theta^k)$
2. prox step. $\theta^{\text{tent}} = \operatorname{argmin}_{\theta} \left(g(\theta) + \frac{1}{2h^k} \|\theta - \theta^{k+1/2}\|^2 \right)$
3. if $F(\theta^{\text{tent}}) \leq F(\theta^k)$,
 - (a) set $\theta^{k+1} = \theta^{\text{tent}}$, $h^{k+1} = 1.2h^k$
 - (b) quit if $\left\| \nabla f(\theta^{k+1}) - \frac{1}{h^k} (\theta^{k+1} - \theta^{k+1/2}) \right\| \leq \epsilon$
4. else set $h^k := 0.5h^k$ and go to step 1

Prox-gradient method convergence

- ▶ prox-gradient method always finds a stationary point
 - ▶ suitably defined for non-differentiable functions
 - ▶ assuming some technical conditions hold
- ▶ for *convex problems*
 - ▶ prox-gradient method is *non-heuristic*
 - ▶ for any starting point θ^1 , $F(\theta^k) \rightarrow F^*$ as $k \rightarrow \infty$
- ▶ for *non-convex problems*
 - ▶ prox-gradient method is *heuristic*
 - ▶ we can (and often do) have $F(\theta^k) \not\rightarrow F^*$

Prox-gradient for regularized ERM

Prox-gradient for sum squares regularizer

- ▶ let's apply prox-gradient method to $F(\theta) = \mathcal{L}(\theta) + \lambda \|\theta\|_2^2$
 - ▶ $f(\theta) = \mathcal{L}(\theta)$
 - ▶ $g(\theta) = \lambda \|\theta\|_2^2 = \lambda \theta_1^2 + \dots + \lambda \theta_d^2$
- ▶ in prox step, we need to minimize $\lambda \theta_i^2 + \frac{1}{2h^k}(\theta_i - \theta_i^{k+1/2})^2$ over θ_i
- ▶ solution is $\theta_i = \frac{1}{1+2\lambda h^k} \theta_i^{k+1/2}$
- ▶ so prox step just shrinks the gradient step $\theta^{k+1/2}$ by the factor $\frac{1}{1+2\lambda h^k}$
- ▶ prox-gradient iteration:
 1. gradient step: $\theta^{k+1/2} = \theta^k - h^k \nabla \mathcal{L}(\theta^k)$
 2. prox step: $\theta^{k+1} = \frac{1}{1+2\lambda h^k} \theta^{k+1/2}$

Prox-gradient for ℓ_1 regularizer

- ▶ let's apply prox-gradient method to $F(\theta) = \mathcal{L}(\theta) + \lambda \|\theta\|_1$
 - ▶ $f(\theta) = \mathcal{L}(\theta)$
 - ▶ $g(\theta) = \lambda \|\theta\|_1 = \lambda |\theta_1| + \dots + \lambda |\theta_d|$
- ▶ in prox step, we need to minimize $\lambda |\theta_i| + \frac{1}{2h^k} (\theta_i - \theta_i^{k+1/2})^2$ over θ_i

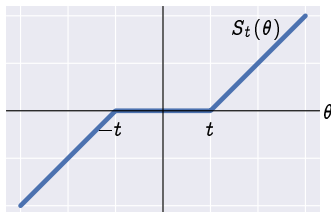
- ▶ solution is

$$\theta_i^{k+1} = \begin{cases} \theta_i^{k+1/2} - \lambda h^k & \theta_i^{k+1/2} > \lambda h^k \\ 0 & |\theta_i^{k+1/2}| \leq \lambda h^k \\ \theta_i^{k+1/2} + \lambda h^k & \theta_i^{k+1/2} < -\lambda h^k \end{cases}$$

- ▶ called *soft threshold function*
- ▶ sometimes written as

$$\begin{aligned} \theta_i^{k+1} &= S_{\lambda h^k}(\theta_i^{k+1/2}) = \text{sign}(\theta_i^{k+1/2}) (|\theta_i^{k+1/2}| - \lambda h^k)_+ \\ &= (\theta_i^{k+1/2} - \lambda h^k)_+ - (-\theta_i^{k+1/2} - \lambda h^k)_+ \end{aligned}$$

Soft threshold function



► prox-gradient iteration for regularized ERM with ℓ_1 regularization:

1. gradient step: $\theta^{k+1/2} = \theta^k - h^k \nabla \mathcal{L}(\theta^k)$

2. prox step: $\theta_i^{k+1} = S_{\lambda h^k}(\theta_i^{k+1/2})$ for $i = 1, \dots, d$.

- the soft threshold step shrinks all coefficients
- and sets the small ones to zero

Prox-gradient step for nonnegative regularizer

- ▶ let's apply prox-gradient method to $F(\theta) = \mathcal{L}(\theta) + r(\theta)$, where $r(\theta) = 0$ for $\theta \geq 0$, ∞ otherwise
 - ▶ $f(\theta) = \mathcal{L}(\theta)$
 - ▶ $g(\theta) = q(\theta_1) + \dots + q(\theta_d)$
- ▶ in prox step, we need to minimize $q(\theta_i) + \frac{1}{2h^k}(\theta_i - \theta_i^{k+1/2})^2$ over θ_i
- ▶ solution is $\theta_i = \left(\theta_i^{k+1/2}\right)_+$
- ▶ so prox step just replaces the gradient step $\theta_i^{k+1/2}$ with its positive part
- ▶ prox gradient iteration:
 1. gradient step: $\theta^{k+1/2} = \theta^k - h^k \nabla \mathcal{L}(\theta^k)$
 2. prox step: $\theta^{k+1} = \left(\theta^{k+1/2}\right)_+$