

# A short course on distributed optimization for machine learning

Jong-Han Kim

EIC7024  
Kyung Hee University

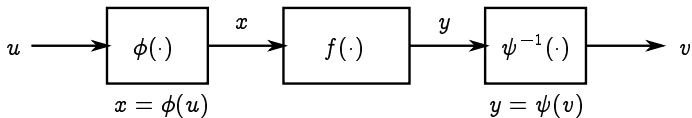
## Course contents

- ▶ Topics
  - ▶ Machine learning problems
    - ▶ Regression
    - ▶ Classification
  - ▶ Optimization for machine learning
    - ▶ Optimization algorithms
    - ▶ Distributed optimization and statistical learning
- ▶ Schedules
  - ▶ Wk.1: supervised learning, empirical risk minimization, loss functions
  - ▶ Wk.2: regularization, classification problems
  - ▶ Wk.3: optimization
  - ▶ Wk.4: distributed optimization

# Predictors

## Information flow

*raw data*                      *feature*                      *outcome*                      *raw output*



student

hrs. sleep  
hrs. study  
hrs. facebook  
alcohol freq.  
⋮

exam score

grade

## Data fitting

- ▶ we think  $y \in \mathbf{R}$  and  $x \in \mathbf{R}^d$  are (approximately) related by

$$y \approx f(x)$$

- ▶  $x$  is called the *independent variable* or *feature vector*
- ▶  $y$  is called the *outcome* or *response* or *target* or *label* or *dependent variable*
- ▶ often  $y$  is something we want to predict
- ▶ we don't know the 'true' relationship between  $x$  and  $y$

## Features

often  $x$  is a vector of features:

- ▶ documents
  - ▶  $x$  is word count histogram for a document
- ▶ patient data
  - ▶  $x$  are patient attributes, test results, symptoms
- ▶ customers
  - ▶  $x$  is purchase history and other attributes of a customer

## Where features come from

- ▶ we use  $u$  to denote the raw input data, such as a vector, word or text, image, video, audio, ...
- ▶  $x = \phi(u)$  is the corresponding *feature vector*
- ▶ the function  $\phi$  is called the *embedding* or *feature function*
- ▶  $\phi$  might be very simple or quite complicated
- ▶ similarly, the raw output data  $v$  can be featurized as  $y = \psi(v)$
- ▶ often we take  $\phi(u)_1 = x_1 = 1$ , the *constant feature*
- ▶ (much more on these ideas later)

## Data and prior knowledge

- ▶ we are given data  $x^1, \dots, x^n \in \mathbf{R}^d$  and  $y^1, \dots, y^n \in \mathbf{R}$
- ▶  $(x^i, y^i)$  is the  $i$ th *data pair* or *observation* or *example*
- ▶ we also (might) have *prior knowledge* about what  $f$  might look like
  - ▶ e.g.,  $f$  is smooth or continuous:  $f(x) \approx f(\tilde{x})$  when  $x$  is near  $\tilde{x}$
  - ▶ or we might know  $y \geq 0$



## Predictor

- ▶ we seek a *predictor* or *model*  $g : \mathbf{R}^d \rightarrow \mathbf{R}$
- ▶ for feature vector  $x$ , our prediction (of  $y$ ) is  $\hat{y} = g(x)$
- ▶ predictor  $g$  is chosen based on both data and prior knowledge
- ▶ in terms of raw data, our predictor is

$$\hat{v} = \psi^{-1}(g(\phi(u)))$$

(with a slight variation when  $\psi$  is not invertible)

- ▶  $\hat{y}^i \approx y^i$  means our predictor does well on  $i$ th data pair
- ▶ *but our real goal is to have  $\hat{y} \approx y$  for  $(x, y)$  pairs we have not seen*

## Prediction methods

- ▶ fraud, psychic powers, telepathy, magic sticks, incantations, crystals, hunches, statistics, AI, machine learning, data science
- ▶ and many algorithms . . .
- ▶ example: nearest neighbor predictor
  - ▶ given  $x$ , find its nearest neighbor  $x^i$  among given data
  - ▶ then predict  $\hat{y} = g(x) = y^i$

A learning algorithm is a recipe for producing a predictor given data

## Linear predictors

## Linear predictor

- ▶ predictors that are linear functions of  $x$  are widely used
- ▶ a linear predictor has the form

$$g(x) = \theta^\top x$$

for some vector  $\theta \in \mathbf{R}^d$ , called the *predictor parameter vector*

- ▶ also called a *regression model*
- ▶  $x_j$  is the  $j$ th feature, so the prediction is a linear combination of features

$$\hat{y} = g(x) = \theta_1 x_1 + \cdots + \theta_d x_d$$

- ▶ we get to choose the predictor parameter vector  $\theta \in \mathbf{R}^d$
- ▶ sometimes we write  $g_\theta(x)$  to emphasize the dependence on  $\theta$

## Interpreting a linear predictor

$$\hat{y} = g(x) = \theta_1 x_1 + \cdots + \theta_d x_d$$

- ▶  $\theta_3$  is the amount that prediction  $\hat{y} = g(x)$  increases when  $x_3$  increases by 1
  - ▶ particularly interpretable when  $x_3$  is Boolean (only takes values 0 or 1)
- ▶  $\theta_7 = 0$  means that the prediction does not depend on  $x_7$
- ▶  $\theta$  small means predictor is insensitive to changes in  $x$ :

$$|g(x) - g(\tilde{x})| = |\theta^\top x - \theta^\top \tilde{x}| = |\theta^\top (x - \tilde{x})| \leq \|\theta\| \|x - \tilde{x}\|$$

## Affine predictor

- ▶ suppose the first feature is constant,  $x_1 = 1$
- ▶ the linear predictor  $g$  is then an *affine function* of  $x_{2:d}$ , i.e., linear plus a constant

$$g(x) = \theta^\top x = \theta_1 + \theta_2 x_2 + \cdots + \theta_d x_d$$

- ▶  $\theta_1$  is called the *offset* or *constant term* in the predictor
- ▶  $\theta_1$  is the prediction when all features (except the constant) are zero

# Empirical risk minimization

## Loss function

a *loss* or *risk* function  $\ell : \mathbf{R} \times \mathbf{R} \rightarrow \mathbf{R}$  quantifies how well (more accurately, how badly)  $\hat{y}$  approximates  $y$

- ▶ smaller values of  $\ell(\hat{y}, y)$  indicate that  $\hat{y}$  is a good approximation of  $y$
- ▶ typically  $\ell(y, y) = 0$  and  $\ell(\hat{y}, y) \geq 0$  for all  $\hat{y}, y$

### examples

- ▶ *quadratic loss*:  $\ell(\hat{y}, y) = (\hat{y} - y)^2$
- ▶ *absolute loss*:  $\ell(\hat{y}, y) = |\hat{y} - y|$



## Empirical risk

how well does the predictor  $g$  fit a data set  $(x^i, y^i)$ ,  $i = 1, \dots, n$ , with loss  $\ell$ ?

- ▶ the *empirical risk* is the average loss over the data points,

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n \ell(\hat{y}^i, y^i) = \frac{1}{n} \sum_{i=1}^n \ell(g(x^i), y^i)$$

- ▶ if  $\mathcal{L}$  is small, the predictor predicts the given data well
- ▶ when the predictor is parametrized by  $\theta$ , we write

$$\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(g_{\theta}(x^i), y^i)$$

to show the dependence on the predictor parameter  $\theta$

## Empirical risk minimization

- ▶ choosing the parameter  $\theta$  in a parametrized predictor  $g_\theta(x)$  is called *fitting the predictor* (to data)
- ▶ *empirical risk minimization (ERM)* is a general method for fitting a parametrized predictor
- ▶ ERM: *choose  $\theta$  to minimize empirical risk  $\mathcal{L}(\theta)$*
- ▶ thus, ERM chooses  $\theta$  by attempting to match given data
- ▶ often there is no analytic solution to this minimization problem, so we use *numerical optimization* to find  $\theta$  that minimizes  $\mathcal{L}(\theta)$  (more on this topic later)

## Least squares linear regression

## Least squares linear regression

- ▶ linear predictor  $\hat{y} = g_{\theta}(x) = \theta^{\top} x$
- ▶  $\theta \in \mathbf{R}^d$  is the model parameter
- ▶ we'll use square loss function  $\ell(\hat{y}, y) = (\hat{y} - y)^2$
- ▶ empirical risk is MSE

$$\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n (\theta^{\top} x^i - y^i)^2$$

- ▶ ERM: choose model parameter  $\theta$  to minimize MSE
- ▶ called *linear least squares fitting* or *linear regression*

## Least squares formulation

- express MSE in matrix notation as

$$\begin{aligned}\mathcal{L}(\theta) &= \frac{1}{n} \sum_{i=1}^n (\theta^\top x^i - y^i)^2 = \frac{1}{n} \{ (\theta^\top x^1 - y^1)^2 + \dots + (\theta^\top x^n - y^n)^2 \} \\ &= \frac{1}{n} \left\| \begin{bmatrix} (x^1)^\top \theta - y^1 \\ \vdots \\ (x^n)^\top \theta - y^n \end{bmatrix} \right\|^2 \\ &= \frac{1}{n} \left\| \underbrace{\begin{bmatrix} (x^1)^\top \\ \vdots \\ (x^n)^\top \end{bmatrix}}_X \theta - \underbrace{\begin{bmatrix} y^1 \\ \vdots \\ y^n \end{bmatrix}}_y \right\|^2 = \frac{1}{n} \|X\theta - y\|^2\end{aligned}$$

- ERM is a *least squares problem*: choose  $\theta$  to minimize  $\|X\theta - y\|^2$   
(factor  $1/n$  doesn't affect choice of  $\theta$ )

## Least squares solution

- ▶ assuming  $X$  has linearly independent columns (which implies  $n \geq d$ ), there is a unique optimal  $\theta$

$$\theta^* = (X^T X)^{-1} X^T y = X^\dagger y$$

- ▶ standard algorithm:
  - ▶ compute QR factorization  $X = QR$  (e.g., Gram-Schmidt) (with orthogonal  $Q$  and invertible upper triangular  $R$ )
  - ▶ compute  $Q^T y$
  - ▶ solve  $R\theta^* = Q^T y$  by back substitution
- ▶ in Julia: `theta_opt = X \ y`
- ▶ in Python: `theta_opt = np.linalg.lstsq(X,y,rcond=None)[0]`
- ▶ complexity is  $2d^2n$  flops

## Data matrix

- ▶ the  $n \times d$  matrix

$$X = \begin{bmatrix} (x^1)^\top \\ \vdots \\ (x^n)^\top \end{bmatrix}$$

is called the *data matrix*

- ▶  $i$ th row of  $X$  is  $i$ th feature vector, transposed
- ▶  $j$ th column of  $X$  gives values of  $j$ th feature  $x_j$  across our data set
- ▶  $X_{ij}$  is the value of  $j$ th feature for the  $i$ th data point

## Constant fit

- ▶ the simplest feature vector is constant:  $x = \phi(u) = 1$   
(doesn't depend on  $u$ !)
- ▶ corresponding predictor is a constant function:  $g(x) = \theta_1$
- ▶ data matrix is  $X = \mathbf{1}_n$
- ▶ so  $X^\dagger = (X^\top X)^{-1} X^\top = (1/n) \mathbf{1}^\top$  and

$$\theta^* = X^\dagger y = \mathbf{1}^\top y / n = \text{avg}(y)$$

- ▶ *the average of the outcome values is the best constant predictor* (for square loss)
- ▶ optimal RMSE is standard deviation of outcome values

$$\left( \frac{1}{n} \sum_{i=1}^n (\text{avg}(y) - y^i)^2 \right)^{1/2}$$



## Regression

- ▶ with  $u \in \mathbf{R}^{d-1}$ :  $x = \phi(u) = (1, u)$
- ▶ same as  $x_1 = 1$  (the first feature is constant)
- ▶ predictor has form

$$\hat{y} = \theta^\top x = \theta_1 + \theta_{2:d}^\top u$$

an affine function of  $u$

## Straight line fit

- ▶ with  $u \in \mathbf{R}$ ,  $x = (1, u) \in \mathbf{R}^2$
- ▶ model is  $\hat{y} = g(x) = \theta_1 + \theta_2 u$
- ▶ this model is called *straight-line fit*
- ▶ when  $u$  is time, it's called the *trend line*
- ▶ when  $u$  is the whole market return, and  $y$  is an asset return,  $\theta_2$  is called ' $\beta$ '

## Constant versus straight-line fit models

- ▶ for the constant model, we choose  $\theta_1$  to minimize

$$\frac{1}{n} \sum_{i=1}^n (\theta_1 - y^i)^2$$

- ▶ for the straight-line model, we choose  $(\theta_1, \theta_2)$  to minimize

$$\frac{1}{n} \sum_{i=1}^n (\theta_1 + \theta_2 u^i - y^i)^2$$

- ▶ for optimal choices, this value is less than or equal to the one above (since we can take  $\theta_2 = 0$  in the straight-line model)
- ▶ so the RMS error of the straight-line fit is no more than the standard deviation

## Penalty functions and error histograms


## Loss and penalty functions

- ▶ empirical risk (or average loss) is  $\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(\theta^\top x^i, y^i)$
- ▶ the loss function  $\ell(\hat{y}, y)$  penalizes deviation between the predicted value  $\hat{y}$  and the observed value  $y$
- ▶ common form for loss function:  $\ell(\hat{y}, y) = p(\hat{y} - y)$
- ▶  $p$  is the *penalty function*
- ▶ e.g., the square penalty  $p^{\text{sqr}}(r) = r^2$
- ▶  $r = \hat{y} - y$  is the *prediction error* or *residual*

## Penalty functions

- ▶ the penalty function tells us how much we object to different values of prediction error
- ▶ usually  $p(0) = 0$  and  $p(r) \geq 0$  for all  $r$
- ▶ if  $p$  is *symmetric*, i.e.,  $p(-r) = p(r)$ , we care only about the magnitude (absolute value) of prediction error
- ▶ if  $p$  is *asymmetric*, i.e.,  $p(-r) \neq p(r)$ , it bothers us more to over- or underestimate

## Square versus absolute value penalty



- ▶ for square penalty  $p^{\text{sqr}}(r) = r^2$ 
  - ▶ for small prediction errors, penalty is very small (small squared)
  - ▶ for large prediction errors, penalty is very large (large squared)
- ▶ for absolute penalty  $p^{\text{abs}}(r) = |r|$ 
  - ▶ for small prediction errors, penalty is large (compared to square)
  - ▶ for large prediction errors, penalty is small (compared to square)

## Predictors and choice of penalty function


- ▶ choice of penalty function depends on how you feel about large, small, positive, or negative prediction errors
- ▶ different choices of penalty function yield different predictor parameters
- ▶ choice of penalty function *shapes* the histogram of prediction errors, *i.e.*,

$$r^1, \dots, r^n$$

(usually divided into bins and displayed as bar graph distribution)



## Histogram of residuals



- ▶ artificial data with  $n = 300$  and  $d = 30$ , using 50/50 test/train split
- ▶ plots show histogram of residuals  $r^1, \dots, r^n$
- ▶ tilted loss results in distribution with most residuals  $r^i < 0$ , i.e., predictor prefers  $\hat{y}^i < y^i$

# Robust fitting


# Outliers

- ▶ in some applications, a few data points are 'way off', or just 'wrong'
- ▶ occurs due to transcription errors, error in decimal point position, *etc.*
- ▶ these points are called *outliers*
- ▶ even a few outliers in a data set can result in a poor predictor
- ▶ several standard methods are used to remove outliers, or reduce their impact
- ▶ one simple method:
  - ▶ create predictor from data set
  - ▶ flag data points with large prediction errors as outliers
  - ▶ remove them from the data set and repeat

## Robust penalty functions

- ▶ we say a penalty function is *robust* if it has low sensitivity to outliers
- ▶ robust penalty functions grow more slowly for large prediction error values than the square penalty
- ▶ and so 'allow' the predictor to have a few large prediction errors (presumably for the outliers)
- ▶ so they handle outliers more gracefully
- ▶ a *robust predictor* might fit, e.g., 98% of the data very well

## Huber loss





- ▶ the *Huber* penalty function is

$$p^{\text{hub}}(r) = \begin{cases} r^2 & \text{if } |r| \leq \alpha \\ \alpha(2|r| - \alpha) & \text{if } |r| > \alpha \end{cases}$$


- ▶  $\alpha$  is a parameter
- ▶ quadratic for small  $r$ , affine for large  $r$

## Huber loss

- ▶ linear growth for large  $r$  makes fit less sensitive to outliers
- ▶ ERM with Huber loss is called a *robust* prediction method



## Log Huber




- quadratic for small  $y$ , logarithmic for large  $y$

$$p^{\text{dh}}(y) = \begin{cases} y^2 & \text{if } |y| \leq \alpha \\ \alpha^2(1 - 2 \log(\alpha) + \log(y^2)) & \text{if } |y| > \alpha \end{cases}$$

- diminishing incremental penalty at large  $y$


## Log Huber



- even less sensitive to outliers than Huber



# Error distribution



# Regularization

## Sensitivity

- ▶ we have a linear predictor  $\hat{y} = g(x) = \theta^\top x$
- ▶ if  $|\theta_i|$  is large, then the prediction is very sensitive to  $x_i$   
(i.e., small changes in  $x_i$  lead to large changes in the prediction)
- ▶ large sensitivity can lead to overfit, poor generalization  
(which would turn up in validation)
- ▶ for  $x_1 = 1$  (the constant feature), there is no sensitivity, since the feature does not change
- ▶ suggests that we would like  $\theta$  (or  $\theta_{2:d}$  if  $x_1 = 1$ ) not too large

## Regularizer

- ▶ we will measure the size of  $\theta$  using a *regularizer* function  $r : \mathbb{R}^d \rightarrow \mathbb{R}$
- ▶  $r(\theta)$  is a measure of the size of  $\theta$  (or  $\theta_{2:d}$ )

- ▶ *quadratic regularizer* (a.k.a.  $\ell_2$  or sum-of-squares):

$$r(\theta) = \|\theta\|^2 = \theta_1^2 + \dots + \theta_d^2$$

- ▶ *absolute value regularizer* (a.k.a.  $\ell_1$ ):

$$r(\theta) = \|\theta\|_1 = |\theta_1| + \dots + |\theta_d|$$

## Regularized empirical risk minimization

- ▶ predictor should fit the given data well, *i.e.*, we want empirical risk

$$\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(\theta^\top x^i, y^i)$$

to be small

- ▶ predictor should not be too sensitive, *i.e.*, we want  $r(\theta)$  small
- ▶ to trade off these two objectives, form *regularized empirical risk*

$$\mathcal{L}(\theta) + \lambda r(\theta)$$

where  $\lambda \geq 0$  is the *regularization parameter* (or *hyper-parameter*)

- ▶ *regularized empirical risk minimization* (RERM): choose  $\theta$  to minimize regularized empirical risk
- ▶ an optimization problem

## Regularized empirical risk minimization

- ▶ for  $\lambda = 0$ , RERM reduces to ERM
- ▶ RERM produces a *family* of predictors, one for each value of  $\lambda$
- ▶ in practice, we choose a few tens of values of  $\lambda$ , usually logarithmically spaced over a wide range
- ▶ use validation to choose among the candidate predictors
- ▶ we choose the largest value of  $\lambda$  that gives near minimum test error (*i.e.*, least sensitive predictor that generalizes well)

## Ridge regression

- ▶ *ridge regression*: square loss and regularizer  $r(\theta) = \|\theta\|^2$  (or  $\|\theta_{2:d}\|^2$  if  $x_1 = 1$ )
- ▶ also called *Tykhonov regularized least squares*
- ▶ regularized empirical risk is

$$\begin{aligned}\mathcal{L}(\theta) + \lambda r(\theta) &= \|X\theta - y\|^2 + \lambda \|\theta\|^2 \\ &= \left\| \begin{bmatrix} X \\ \sqrt{\lambda}I \end{bmatrix} \theta - \begin{bmatrix} y \\ 0 \end{bmatrix} \right\|^2\end{aligned}$$

- ▶ so optimal  $\theta$  is

$$\theta^* = \begin{bmatrix} X \\ \sqrt{\lambda}I \end{bmatrix}^\dagger \begin{bmatrix} y \\ 0 \end{bmatrix} = (X^T X + \lambda I)^{-1} X^T y$$

- ▶ (how do you modify this to handle  $r(\theta) = \|\theta_{2:d}\|^2$ ?)

# Regularizers



## Regularizers

- ▶ motivation:
  - ▶ large  $\theta_i$  makes prediction  $\theta^T x$  sensitive to value of  $x_i$
  - ▶ so we want  $\theta$  (or  $\theta_{2:d}$  if  $x_1 = 1$ ) small
- ▶ regularizer  $r : \mathbf{R}^d \rightarrow \mathbf{R}$  measures the size of  $\theta$
- ▶ usually regularizer is *separable*,

$$r(\theta) = q(\theta_1) + \cdots + q(\theta_d)$$

where  $q : \mathbf{R} \rightarrow \mathbf{R}$  is a penalty function for the predictor coefficients

## Sum squares regularizer

- ▶ *sum squares* regularizer uses square penalty  $q^{\text{sqr}}(a) = a^2$

$$r(\theta) = \|\theta\|^2 = \theta_1^2 + \cdots + \theta_d^2$$

- ▶ also called *quadratic*, *Tychonov*, or  $\ell_2$  regularizer

## $\ell_1$ regularizer

- ▶ *sum absolute* or  $\ell_1$  regularizer uses absolute value penalty  $q^{\text{abs}}(a) = |a|$

$$r(\theta) = \|\theta\|_1 = |\theta_1| + \cdots + |\theta_d|$$

- ▶  $\|\theta\|_1$  is  $\ell_1$  *norm* of  $\theta$
- ▶ like the Euclidean or  $\ell_2$  norm  $\|\theta\|$ , it is a norm, *i.e.*, a measure of the size of the vector  $\theta$
- ▶ Euclidean norm is often written as  $\|\theta\|_2$  to distinguish it from the  $\ell_1$  norm
- ▶ they are both members of the *p-norm family*, defined as

$$\|\theta\|_p = (|\theta_1|^p + \cdots + |\theta_d|^p)^{1/p}$$

for  $p \geq 1$

## Lasso regression

- ▶ use square loss  $\ell(\hat{y}, y) = (\hat{y} - y)^2$
- ▶ choosing  $\theta$  to minimize  $\mathcal{L}(\theta) + \lambda \|\theta\|_2^2$  is called *ridge regression*
- ▶ choosing  $\theta$  to minimize  $\mathcal{L}(\theta) + \lambda \|\theta\|_1$  is called *lasso regression*
- ▶ invented by (Stanford's) Rob Tibshirani, 1994
- ▶ widely used in advanced machine learning
- ▶ unlike ridge regression, there is no formula for the lasso parameter vector
- ▶ but we can efficiently compute it anyway (since it's convex)
- ▶ the lasso regression model has some interesting properties

## Sparsifying regularizers

## Sparse coefficient vector

- ▶ suppose  $\theta$  is sparse, *i.e.*, many of its entries are zero
- ▶ prediction  $\theta^T x$  does not depend on features  $x_i$  for which  $\theta_i = 0$
- ▶ this means we select *some* features to use (*i.e.*, those with  $\theta_i \neq 0$ )
- ▶ (possible) practical benefits of sparse  $\theta$ :
  - ▶ can improve performance when many features are actually irrelevant
  - ▶ makes predictor *simpler to interpret*

## Sparse coefficient vectors via $\ell_1$ regularization

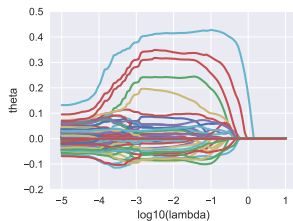
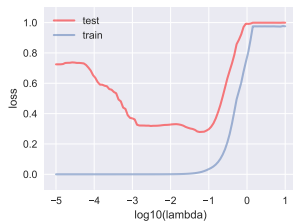
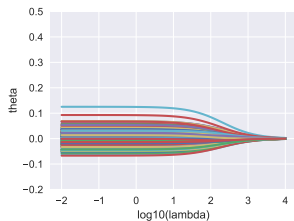
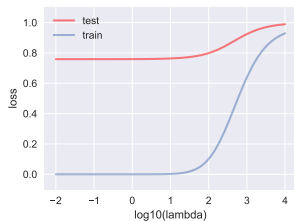
*using  $\ell_1$  regularization leads to sparse coefficient vectors*

$r(\theta) = \|\theta\|_1$  is called a *sparsifying regularizer*

rough explanation:

- ▶ for square penalty, once  $\theta_i$  is small,  $\theta_i^2$  is very small
- ▶ so incentive for sum squares regularizer to make a coefficient smaller decreases once it is small
- ▶ for absolute penalty, incentive to make  $\theta_i$  smaller keeps up all the way until it's zero

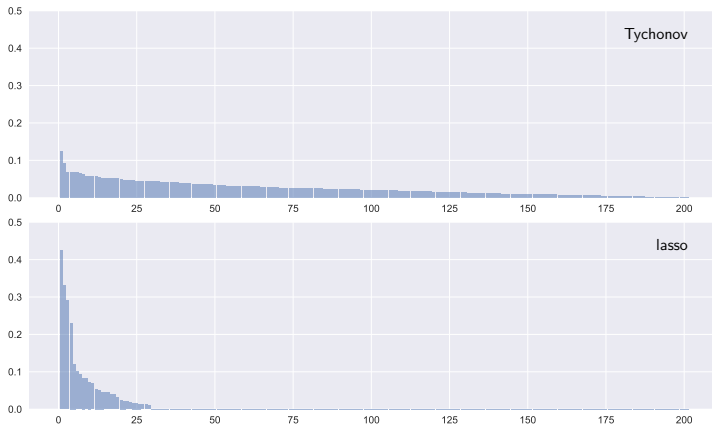
## Example



- ▶ artificially generated 50 data points, 200 features
- ▶ only a few features are relevant
- ▶ left hand plots use Tychonov, right hand use lasso

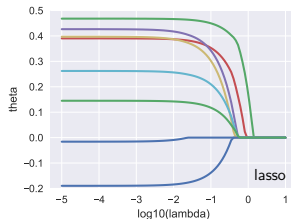
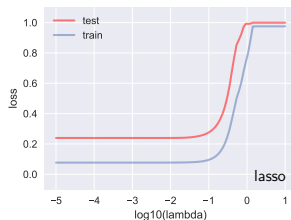
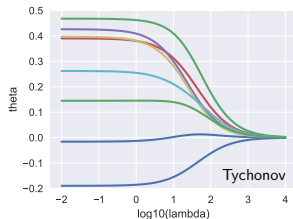
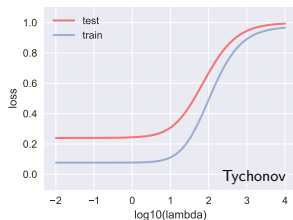


## Example



- ▶ sorted  $|\theta_i|$  at optimal  $\lambda$
- ▶ lasso solution has only 35 non-zero components

## Example



- choose  $\lambda$  based on regularization path with test data
- keep features corresponding to largest components of  $\theta$  and *retrain*
- plots above use most important 7 features identified by lasso

## Even stronger sparsifiers

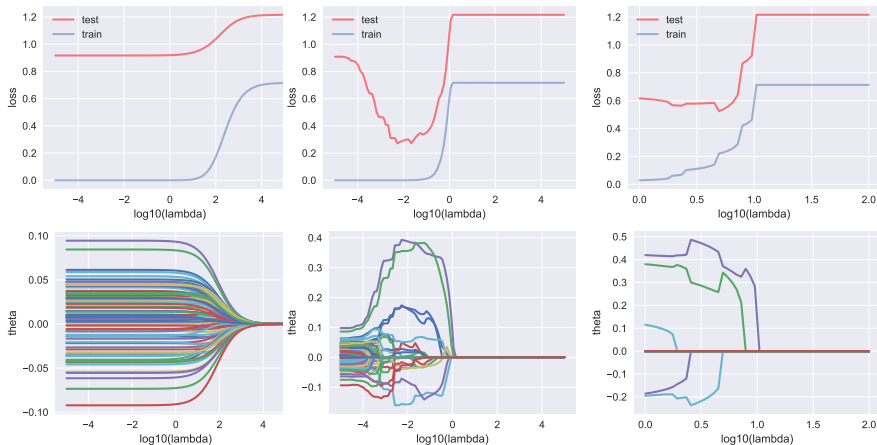
- ▶  $q(a) = |a|^{1/2}$
- ▶ called  $\ell_{0.5}$  regularizer
- ▶ but you shouldn't use this term since

$$\left(|\theta_1|^{0.5} + \dots + |\theta_d|^{0.5}\right)^2$$

is not a norm (see VMLS)

- ▶ 'stronger' sparsifier than  $\ell_1$
- ▶ but not convex so computing  $\theta$  is heuristic

## Example



►  $\ell_2$ ,  $\ell_1$ , and square root regularization

## Nonnegative regularizer

## Nonnegative coefficients

- ▶ in some cases we know or require that  $\theta_i \geq 0$
- ▶ this means that when  $x_i$  increases, so must our prediction
- ▶ we can think of this constraint as regularization with penalty function

$$q(a) = \begin{cases} 0 & a \geq 0 \\ \infty & a < 0 \end{cases}$$

- ▶ example:  $y$  is lifespan,  $x_i$  measures healthy behavior  $i$
- ▶ with quadratic loss, called *nonnegative least squares* (NNLS)
- ▶ common heuristic for nonnegative least squares: use  $(\theta^{\text{ls}})_+$  (works poorly)

# Classification

## Boolean classification

- ▶ supervised learning is called *boolean classification* when raw output variable  $v$  is a categorical that can take two possible values
- ▶ we denote these  $-1$  and  $1$ , and they often correspond to  $\{\text{FALSE}, \text{TRUE}\}$  or  $\{\text{NEGATIVE}, \text{POSITIVE}\}$
- ▶ for a data record  $u^i, v^i$ , the value  $v^i \in \{-1, 1\}$  is called the *class* or *label*
- ▶ a *boolean classifier* predicts label  $\hat{v}$  given raw input  $u$

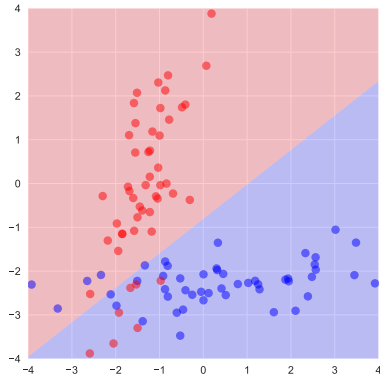


## Classification



- ▶ here  $u \in \mathbb{R}^2$
- ▶ red points have  $v^i = -1$ , blue points have  $v^i = 1$
- ▶ we'd like a predictor that maps any  $u \in \mathbb{R}^2$  into prediction  $\hat{v} \in \{-1, 1\}$

## Example: Least squares classifier



- ▶ embed  $x = (1, u)$  and  $y = v$ , apply least squares regression
- ▶ gives  $\hat{y} = \theta_1 + \theta_2 u_1 + \theta_3 u_2$
- ▶ predict using  $\hat{v} = \text{sign}(\hat{y})$
- ▶ 11% of points misclassified at training

# Confusion matrix

## The two types of errors

- ▶ measure performance of a specific predictor on a set of  $n$  data records
- ▶ each data point  $i$  has  $v^i \in \{-1, 1\}$
- ▶ and corresponding prediction  $\hat{v}^i = g(v^i) \in \{-1, 1\}$
- ▶ only four possible values for the data pair  $\hat{v}^i, v^i$ :
  - ▶ *true positive* if  $\hat{v}^i = 1$  and  $v^i = 1$
  - ▶ *true negative* if  $\hat{v}^i = -1$  and  $v^i = -1$
  - ▶ *false negative* or *type II error* if  $\hat{v}^i = -1$  and  $v^i = 1$
  - ▶ *false positive* or *type I error* if  $\hat{v}^i = 1$  and  $v^i = -1$

## Confusion matrix

- ▶ for a predictor and a data set define the *confusion matrix*

$$C = \begin{bmatrix} \# \text{ true negatives} & \# \text{ false negatives} \\ \# \text{ false positives} & \# \text{ true positives} \end{bmatrix} = \begin{bmatrix} C_{\text{tn}} & C_{\text{fn}} \\ C_{\text{fp}} & C_{\text{tp}} \end{bmatrix}$$

(warning: some people use the transpose of  $C$ )

- ▶  $C_{\text{tn}} + C_{\text{fn}} + C_{\text{fp}} + C_{\text{tp}} = n$  (total number of examples)
- ▶  $N_{\text{n}} = C_{\text{tn}} + C_{\text{fp}}$  is number of negative examples
- ▶  $N_{\text{p}} = C_{\text{fn}} + C_{\text{tp}}$  is number of positive examples
- ▶ diagonal entries give numbers of correct predictions
- ▶ off-diagonal entries give numbers of incorrect predictions of the two types

## Some boolean classification measures

- ▶ confusion matrix  $\begin{bmatrix} C_{tn} & C_{fn} \\ C_{fp} & C_{tp} \end{bmatrix}$
- ▶ the basic error measures:
  - ▶ *false positive rate* is  $C_{fp}/n$
  - ▶ *false negative rate* is  $C_{fn}/n$
  - ▶ *error rate* is  $(C_{fn} + C_{fp})/n$
- ▶ error measures some people use:
  - ▶ *true positive rate* or *sensitivity* or *recall* is  $C_{tp}/N_p$
  - ▶ *false alarm rate* is  $C_{fp}/N_n$
  - ▶ *specificity* or *true negative rate* is  $C_{tn}/N_n$
  - ▶ *precision* is  $C_{tp}/(C_{tp} + C_{fp})$

## Neyman-Pearson error

- ▶ *Neyman-Pearson error* over a data set is  $\kappa C_{\text{fn}}/n + C_{\text{fp}}/n$
  - ▶ a scalarization of our two objectives, false positive and false negative rates
  - ▶  $\kappa$  is how much more false negatives irritate us than false positives
  - ▶ when  $\kappa = 1$ , the Neyman-Pearson error is the *error rate*
- 
- ▶ we'll use the Neyman-Pearson error as our scalarized measure

## ERM for classification tasks



## Embedding

- ▶ we embed raw input and output records as  $x = \phi(u)$  and  $y = \psi(v)$
- ▶  $\phi$  is the feature map
- ▶  $\psi$  is *the identity map*,  $\psi(v) = v$
- ▶ un-embed by  $\hat{v} = \text{sign}(\hat{y})$
- ▶ equivalent to  $\hat{v} = \underset{v \in \{-1, 1\}}{\text{argmin}} |\hat{y} - \psi(v)|$
- ▶ i.e., choose the nearest boolean value to the (real) prediction  $\hat{y}$

## ERM

- ▶ given loss function  $\ell(\hat{y}, y)$ , *empirical risk* on a data set is

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n \ell(\hat{y}^i, y^i)$$

- ▶ for linear model  $\hat{y} = \theta^\top x$ , with  $\theta \in \mathbf{R}^d$ ,

$$\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(\theta^\top x^i, y^i)$$

- ▶ ERM: choose  $\theta$  to minimize  $\mathcal{L}(\theta)$
- ▶ regularized ERM: choose  $\theta$  to minimize  $\mathcal{L}(\theta) + \lambda r(\theta)$ , with  $\lambda > 0$

## Loss functions for boolean classification

- ▶ to apply ERM, we need a loss function on embedded variables  $\ell(\hat{y}, y)$
- ▶  $y$  can only take values  $-1$  or  $1$
- ▶ but  $\hat{y} = \theta^\top x \in \mathbf{R}$  can be any real number
- ▶ to specify  $\ell$ , we only need to give two functions (of a scalar  $\hat{y}$ ):
  - ▶  $\ell(\hat{y}, -1)$  is how much  $\hat{y}$  irritates us when  $y = -1$
  - ▶  $\ell(\hat{y}, 1)$  is how much  $\hat{y}$  irritates us when  $y = 1$
- ▶ we can take  $\ell(\hat{y}, 1) = \kappa \ell(-\hat{y}, -1)$ , to reflect that false negatives irritate us a factor  $\kappa$  more than false positives

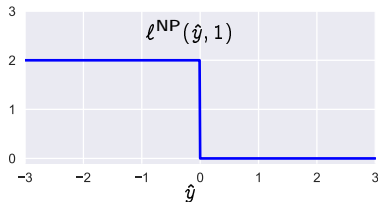
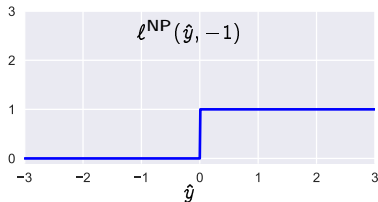
## Neyman-Pearson loss

► Neyman-Pearson loss is

$$\ell^{\text{NP}}(\hat{y}, -1) = \begin{cases} 1 & \hat{y} \geq 0 \\ 0 & \hat{y} < 0 \end{cases}$$

$$\ell^{\text{NP}}(\hat{y}, 1) = \kappa \ell^{\text{NP}}(\hat{y}, -1) = \begin{cases} \kappa & \hat{y} < 0 \\ 0 & \hat{y} \geq 0 \end{cases}$$

► empirical Neyman-Pearson risk  $\mathcal{L}^{\text{NP}}$  is the Neyman-Pearson error





## The problem with Neyman-Pearson loss

- ▶ empirical Neyman-Pearson risk  $\mathcal{L}^{\text{NP}}(\theta)$  is not differentiable, or even continuous (and certainly not convex)
- ▶ worse, its gradient  $\nabla \mathcal{L}^{\text{NP}}(\theta)$  is either zero or undefined
- ▶ so an optimizer does not know how to improve the predictor

## Idea of proxy loss


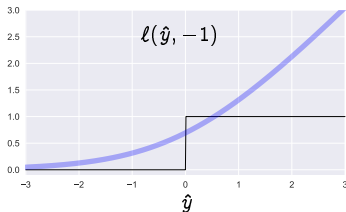
- ▶ we get better results using a *proxy loss* that
  - ▶ approximates, or at least captures the flavor of, the Neyman-Pearson loss
  - ▶ is more easily optimized (e.g., is convex or has nonzero derivative)
- ▶ we want a proxy loss function
  - ▶ with  $\ell(\hat{y}, -1)$  small when  $\hat{y} < 0$ , and larger when  $\hat{y} > 0$
  - ▶ with  $\ell(\hat{y}, +1)$  small when  $\hat{y} > 0$ , and larger when  $\hat{y} < 0$
  - ▶ which has other nice characteristics, e.g., differentiable or convex

## Sigmoid loss



- ▶  $\ell(\hat{y}, -1) = \frac{1}{1 + e^{-\hat{y}}}$ ,  $\ell(\hat{y}, 1) = \kappa \ell(-\hat{y}, -1) = \frac{\kappa}{1 + e^{\hat{y}}}$
- ▶ differentiable approximation of Neyman-Pearson loss
- ▶ but not convex

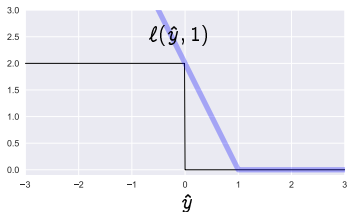
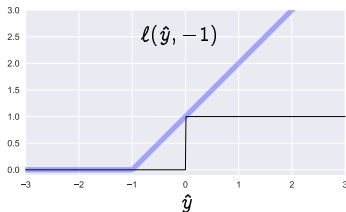
## Logistic loss



- ▶  $\ell(\hat{y}, -1) = \log(1 + e^{\hat{y}})$ ,  $\ell(\hat{y}, 1) = \kappa \ell(-\hat{y}, -1) = \kappa \log(1 + e^{-\hat{y}})$
- ▶ differentiable and convex approximation of Neyman-Pearson loss

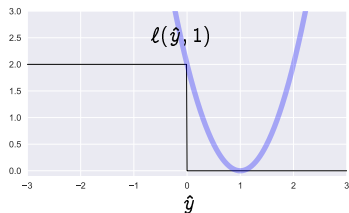
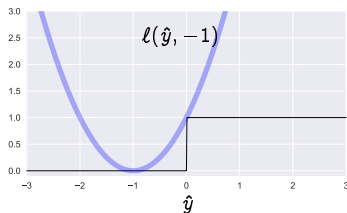


## Hinge loss



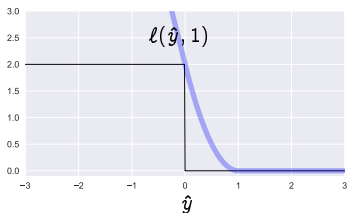
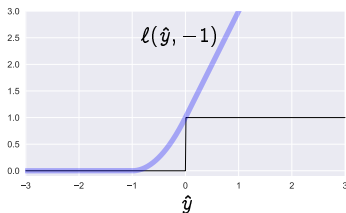
- ▶  $\ell(\hat{y}, -1) = (1 + \hat{y})_+$ ,  $\ell(\hat{y}, 1) = \kappa \ell(-\hat{y}, -1) = \kappa(1 - \hat{y})_+$
- ▶ nondifferentiable but convex approximation of Neyman-Pearson loss

## Square loss



- ▶  $\ell(\hat{y}, -1) = (1 + \hat{y})^2$ ,  $\ell(\hat{y}, 1) = \kappa \ell(-\hat{y}, -1) = \kappa(1 - \hat{y})^2$
- ▶ ERM is least squares problem

## Hubristic loss



- define the *hubristic loss* (huber + logistic) as

$$\ell(\hat{y}, -1) = \begin{cases} 0 & \hat{y} < -1 \\ (\hat{y} + 1)^2 & -1 \leq \hat{y} \leq 0 \\ 1 + 2\hat{y} & \hat{y} > 0 \end{cases}$$

- $\ell(\hat{y}, 1) = \kappa \ell(-\hat{y}, -1)$

## Boolean classifiers

## Least squares classifier

- ▶ use empirical risk with square loss

$$\mathcal{L}(\theta) = \frac{1}{n} \left( \sum_{i: y^i = -1} (1 + \hat{y}^i)^2 + \kappa \sum_{i: y^i = 1} (1 - \hat{y}^i)^2 \right)$$

and your choice of regularizer

- ▶ with sum squares regularizer, this is *least squares classifier*
- ▶ we can minimize  $\mathcal{L}(\theta) + \lambda r(\theta)$  using, e.g., QR factorization

## Logistic regression

- ▶ use empirical risk with logistic loss

$$\mathcal{L}(\theta) = \frac{1}{n} \left( \sum_{i: y^i = -1} \log(1 + e^{\hat{y}^i}) + \kappa \sum_{i: y^i = 1} \log(1 + e^{-\hat{y}^i}) \right)$$

and your choice of regularizer

- ▶ can minimize  $\mathcal{L}(\theta) + \lambda r(\theta)$  using prox-gradient method
- ▶ we will find an actual minimizer if  $r$  is convex

## Support vector machine

(usually abbreviated as *SVM*)

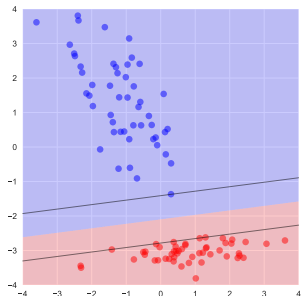
- use empirical risk with hinge loss

$$\mathcal{L}(\theta) = \frac{1}{n} \left( \sum_{i: y^i = -1} (1 + \hat{y}^i)_+ + \kappa \sum_{i: y^i = 1} (1 - \hat{y}^i)_+ \right)$$

and sum squares regularizer

- $\mathcal{L}(\theta) + \lambda r(\theta)$  is convex
- it can be minimized by various methods (but not prox-gradient)

## Support vector machine



- decision boundary is  $\theta^T x = 0$
- black lines show points where  $\theta^T x = \pm 1$
- what is the training risk here?



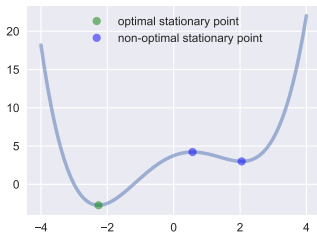
# Optimization problems and algorithms

## Optimization problem

$$\text{minimize } f(\theta)$$

- ▶  $\theta \in \mathbf{R}^d$  is the *variable* or *decision variable*
- ▶  $f : \mathbf{R}^d \rightarrow \mathbf{R}$  is the *objective function*
- ▶ goal is to choose  $\theta$  to minimize  $f$
- ▶  $\theta^*$  is *optimal* means that for all  $\theta$ ,  $f(\theta) \geq f(\theta^*)$
- ▶  $f^* = f(\theta^*)$  is the *optimal value* of the problem
- ▶ optimization problems arise in many fields and applications, including machine learning

## Optimality condition



- ▶ let's assume that  $f$  is *differentiable*, i.e., partial derivatives  $\frac{\partial f(\theta)}{\partial \theta_i}$  exist
- ▶ if  $\theta^*$  is optimal, then  $\nabla f(\theta^*) = 0$
- ▶  $\nabla f(\theta) = 0$  is called the *optimality condition* for the problem
- ▶ there can be points that satisfy  $\nabla f(\theta) = 0$  but are not optimal
- ▶ we call points that satisfy  $\nabla f(\theta) = 0$  *stationary points*
- ▶ not all stationary points are optimal

## Solving optimization problems

- ▶ in some cases, we can solve the problem analytically
- ▶ e.g., least squares: minimize  $f(\theta) = \|X\theta - y\|^2$ 
  - ▶ optimality condition is  $\nabla f(\theta) = 2X^T(X\theta - y) = 0$
  - ▶ this has (unique) solution  $\theta^* = (X^T X)^{-1} X^T y = X^\dagger y$   
(when columns of  $X$  are linearly independent)
- ▶ in other cases, we resort to an *iterative algorithm* that computes a sequence  $\theta^1, \theta^2, \dots$  with, hopefully,  $f(\theta^k) \rightarrow f^*$  as  $k \rightarrow \infty$

## Iterative algorithms

- ▶ *iterative algorithm* computes a sequence  $\theta^1, \theta^2, \dots$
- ▶  $\theta^k$  is called the  $k$ th *iterate*
- ▶  $\theta^1$  is called the *starting point*
- ▶ many iterative algorithms are *descent methods*, which means

$$f(\theta^{k+1}) < f(\theta^k), \quad k = 1, 2, \dots$$

*i.e.*, each iterate is better than the previous one

- ▶ this means that  $f(\theta^k)$  converges, but not necessarily to  $f^*$



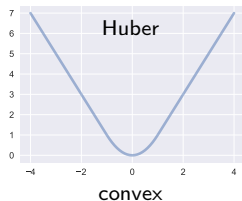
## Stopping criterion

- ▶ in practice, we stop after a finite number  $K$  of steps
- ▶ typical stopping criterion: stop if  $\|\nabla f(\theta^k)\| \leq \epsilon$  or  $k = k^{\max}$
- ▶  $\epsilon$  is a small positive number, the *stopping tolerance*
- ▶  $k^{\max}$  is the maximum number of iterations
- ▶ in words: we stop when  $\theta^k$  is almost a stationary point
- ▶ we hope that  $f(\theta^K)$  is not too much bigger than  $f^*$
- ▶ or more realistically, that  $\theta^K$  is at least useful for our application

## Non-heuristic and heuristic algorithms

- ▶ in some cases we *know* that  $f(\theta^k) \rightarrow f^*$ , for any  $\theta^1$
  - ▶ in words: *we'll get to a solution if we keep iterating*
  - ▶ called *non-heuristic*
- 
- ▶ other algorithms do not guarantee that  $f(\theta^k) \rightarrow f^*$
  - ▶ we can hope that even if  $f(\theta^k) \not\rightarrow f^*$ ,  $\theta^k$  is still useful for our application
  - ▶ called *heuristic*

## Convex functions



- ▶ a function  $f : \mathbf{R}^d \rightarrow \mathbf{R}$  is **convex** if for any  $\theta$ ,  $\tilde{\theta}$ , and  $\alpha$  with  $0 \leq \alpha \leq 1$ ,

$$f(\alpha\theta + (1 - \alpha)\tilde{\theta}) \leq \alpha f(\theta) + (1 - \alpha)f(\tilde{\theta})$$

- ▶ roughly speaking,  $f$  has 'upward curvature'
- ▶ for  $d = 1$ , same as  $f''(\theta) \geq 0$  for all  $\theta$



## Convex optimization

- optimization problem

$$\text{minimize } f(\theta)$$

is called *convex* if the objective function  $f$  is convex

- for convex optimization problem,  $\nabla f(\theta) = 0$  only for  $\theta$  optimal, i.e.,  
*all stationary points are optimal*

- algorithms for convex optimization are non-heuristic
- i.e., *we can solve convex optimization problems* (exactly, in principle)

## Convex ERM problems

- regularized empirical risk function  $f(\theta) = \mathcal{L}(\theta) + \lambda r(\theta)$ , with  $\lambda \geq 0$ ,

$$\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n p(\theta^\top x^i - y^i), \quad r(\theta) = q(\theta_1) + \cdots + q(\theta_d)$$

- $f$  is convex if loss penalty  $p$  and parameter penalty  $q$  functions are convex
- convex penalties: square, absolute, tilted absolute, Huber
- non-convex penalties: log Huber, squareroot

# Gradient method

## Gradient method

- ▶ assume  $f$  is differentiable
- ▶ at iteration  $\theta^k$ , create affine (Taylor) approximation of  $f$  valid near  $\theta^k$

$$\hat{f}(\theta; \theta^k) = f(\theta^k) + \nabla f(\theta^k)^T (\theta - \theta^k)$$

- ▶  $\hat{f}(\theta; \theta^k) \approx f(\theta)$  for  $\theta$  near  $\theta^k$
- ▶ choose  $\theta^{k+1}$  to make  $\hat{f}(\theta^{k+1}; \theta^k)$  small, but with  $\|\theta^{k+1} - \theta^k\|$  not too large
- ▶ choose  $\theta^{k+1}$  to minimize  $\hat{f}(\theta; \theta^k) + \frac{1}{2h^k} \|\theta - \theta^k\|^2$
- ▶  $h^k > 0$  is a *trust parameter* or *step length* or *learning rate*
- ▶ solution is  $\theta^{k+1} = \theta^k - h^k \nabla f(\theta^k)$
- ▶ roughly: take step in direction of negative gradient

## Gradient method update

- choose  $\theta^{k+1}$  to as minimizer of

$$f(\theta^k) + \nabla f(\theta^k)^T(\theta - \theta^k) + \frac{1}{2h^k} \|\theta - \theta^k\|^2$$

- rewrite as

$$f(\theta^k) + \frac{1}{2h^k} \|(\theta - \theta^k) + h^k \nabla f(\theta^k)\|^2 - \frac{h^k}{2} \|\nabla f(\theta^k)\|^2$$

- first and third terms don't depend on  $\theta$
- middle term is minimized (made zero!) by choice

$$\theta = \theta^k - h^k \nabla f(\theta^k)$$

## How to choose step length

- ▶ if  $h^k$  is too large, we can have  $f(\theta^{k+1}) > f(\theta^k)$
- ▶ if  $h^k$  is too small, we have  $f(\theta^{k+1}) < f(\theta^k)$  but progress is slow
- ▶ a simple scheme:
  - ▶ if  $f(\theta^{k+1}) > f(\theta^k)$ , set  $h^{k+1} = h^k/2$ ,  $\theta^{k+1} = \theta^k$  (a *rejected step*)
  - ▶ if  $f(\theta^{k+1}) \leq f(\theta^k)$ , set  $h^{k+1} = 1.2h^k$  (an *accepted step*)
- ▶ reduce step length by half if it's too long; increase it 20% otherwise

## Gradient method summary

choose an initial  $\theta^1 \in \mathbf{R}^d$  and  $h^1 > 0$  (e.g.,  $\theta^1 = 0$ ,  $h^1 = 1$ )

for  $k = 1, 2, \dots, k^{\max}$

1. compute  $\nabla f(\theta^k)$ ; quit if  $\|\nabla f(\theta^k)\|$  is small enough
2. form tentative update  $\theta^{\text{tent}} = \theta^k - h^k \nabla f(\theta^k)$
3. if  $f(\theta^{\text{tent}}) \leq f(\theta^k)$ , set  $\theta^{k+1} = \theta^{\text{tent}}$ ,  $h^{k+1} = 1.2h^k$
4. else set  $h^k := 0.5h^k$  and go to step 2

## Gradient method convergence

- ▶ (assuming some technical conditions hold) we have

$$\|\nabla f(\theta^k)\| \rightarrow 0 \text{ as } k \rightarrow \infty$$

- ▶ i.e., the gradient method always finds a stationary point

- ▶ for *convex problems*

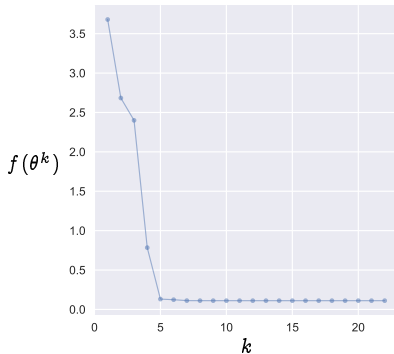
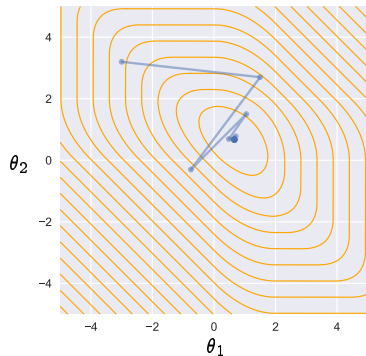
- ▶ gradient method is *non-heuristic*
- ▶ for any starting point  $\theta^1$ ,  $f(\theta^k) \rightarrow f^*$  as  $k \rightarrow \infty$

- ▶ for *non-convex problems*

- ▶ gradient method is *heuristic*
- ▶ we can (and often do) have  $f(\theta^k) \not\rightarrow f^*$


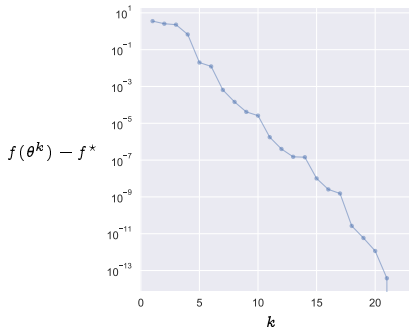


## Example: Convex objective



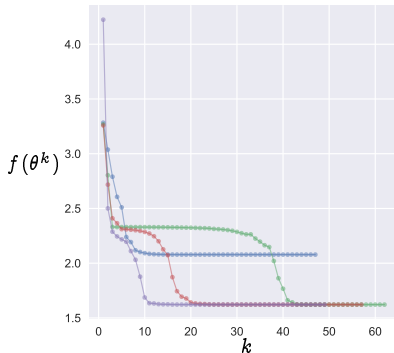
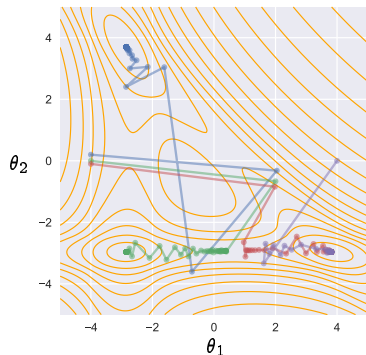
- ▶  $f(\theta) = \frac{1}{3} (p^{\text{hub}}(\theta_1 - 1) + p^{\text{hub}}(\theta_2 - 1) + p^{\text{hub}}(\theta_1 + \theta_2 - 1))$
- ▶  $f$  is convex
- ▶ optimal point is  $\theta^* = (2/3, 2/3)$ , with  $f^* = 1/9$

## Example: Convex objective



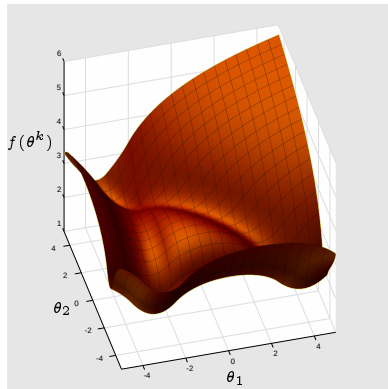

- $f(\theta^k)$  is a decreasing function of  $k$ , (roughly) exponentially
- $\|\nabla f(\theta^k)\| \rightarrow 0$  as  $k \rightarrow \infty$

## Example: Non-convex objective

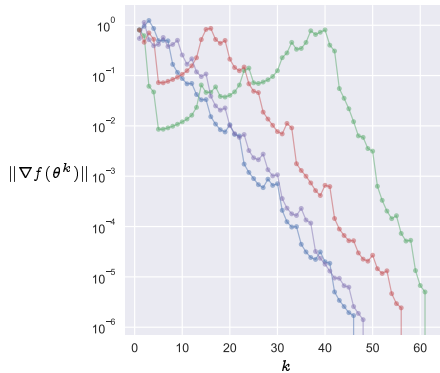
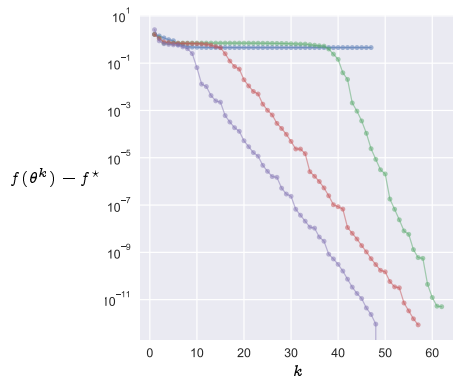


- ▶  $f(\theta) = \frac{1}{3} (p^{\text{lh}}(\theta_1 + 3) + p^{\text{lh}}(2\theta_2 + 6) + p^{\text{lh}}(\theta_1 + \theta_2 - 1))$
- ▶  $f$  is sum of log-Huber functions, so not convex
- ▶ gradient algorithm converges, but limit depends on initial guess

## Example: Non-convex objective



## Example: Non-convex objective



## Gradient method for ERM

## Gradient of empirical risk function

- ▶ empirical risk is sum of terms for each data point

$$\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(\hat{y}^i, y^i) = \frac{1}{n} \sum_{i=1}^n \ell(\theta^T x^i, y^i)$$

- ▶ convex if loss function  $\ell$  is convex in first argument
- ▶ gradient is sum of terms for each data point

$$\nabla \mathcal{L}(\theta) = \nabla \mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n \ell'(\theta^T x^i, y^i) x^i$$

where  $\ell'(\hat{y}, y)$  is derivative of  $\ell$  with respect to its first argument  $\hat{y}$

## Evaluating gradient of empirical risk function

- ▶ compute  $n$ -vector  $\hat{y}^k = X\theta^k$
  - ▶ compute  $n$ -vector  $z^k$ , with entries  $z_i^k = \ell'(\hat{y}_i^k, y^i)$
  - ▶ compute  $d$ -vector  $\nabla \mathcal{L}(\theta^k) = (1/n)X^T z^k$
- 
- ▶ first and third steps are matrix-vector multiplication, each costing  $2nd$  flops
  - ▶ second step costs order  $n$  flops (dominated by other two)
  - ▶ total is  $4nd$  flops



## Prox-gradient method

## Minimizing composite functions

- ▶ want to minimize  $F(\theta) = f(\theta) + g(\theta)$  (called *composite function*)
- ▶  $f$  is differentiable, but  $g$  need not be
- ▶ example: minimize  $\mathcal{L}(\theta) + \lambda r(\theta)$ , with  $r(\theta) = \|\theta\|_1$
- ▶ we'll see idea of gradient method extends directly to composite functions

## Selective linearization

- ▶ at iteration  $k$ , linearize  $f$  *but not*  $g$

$$\hat{F}(\theta; \theta^k) = f(\theta^k) + \nabla f(\theta^k)^T(\theta - \theta^k) + g(\theta)$$

- ▶ want  $\hat{F}(\theta; \theta^k)$  small, but with  $\theta$  near  $\theta^k$
- ▶ choose  $\theta^{k+1}$  to minimize  $\hat{F}(\theta; \theta^k) + \frac{1}{2h^k}\|\theta - \theta^k\|^2$ , with  $h^k > 0$
- ▶ same as minimizing

$$g(\theta) + \frac{1}{2h^k}\|\theta - (\theta^k - h^k \nabla f(\theta^k))\|^2$$

- ▶ for many 'simple' functions  $g$ , this minimization can be done analytically
- ▶ this iteration from  $\theta^k$  to  $\theta^{k+1}$  is called *prox-gradient step*

## Prox-gradient iteration

- prox-gradient iteration has two parts:

1. *gradient step*:  $\theta^{k+1/2} = \theta^k - h^k \nabla f(\theta^k)$

2. *prox step*: choose  $\theta^{k+1}$  to minimize  $g(\theta) + \frac{1}{2h^k} \|\theta - \theta^{k+1/2}\|^2$

( $\theta^{k+1/2}$  is an intermediate iterate, in between  $\theta^k$  and  $\theta^{k+1}$ )

- step 1 handles differentiable part of objective, i.e.,  $f$
- step 2 handles second part of objective, i.e.,  $g$

## Proximal operator

- ▶ given function  $q : \mathbf{R}^d \rightarrow \mathbf{R}$ , and  $\kappa > 0$ ,

$$\mathbf{prox}_{q,\kappa}(v) = \underset{\theta}{\operatorname{argmin}} \left( q(\theta) + \frac{1}{2\kappa} \|\theta - v\|^2 \right)$$

is called the *proximal operator* of  $q$  at  $v$ , with parameter  $\kappa$

- ▶ the prox-gradient step can be expressed as

$$\theta^{k+1} = \mathbf{prox}_{g,h^k}(\theta^{k+1/2}) = \mathbf{prox}_{g,h^k}(\theta^k - h^k \nabla f(\theta^k))$$

- ▶ hence the name prox-gradient iteration

## How to choose step length

- ▶ same as for gradient, but using  $F(\theta) = f(\theta) + g(\theta)$
- ▶ a simple scheme:
  - ▶ if  $F(\theta^{k+1}) > F(\theta^k)$ , set  $h^{k+1} = h^k/2$ ,  $\theta^{k+1} = \theta^k$  (a *rejected step*)
  - ▶ if  $F(\theta^{k+1}) \leq F(\theta^k)$ , set  $h^{k+1} = 1.2h^k$  (an *accepted step*)
- ▶ reduce step length by half if it's too long; increase it 20% otherwise

## Stopping criterion

- ▶ stopping condition for prox-gradient method:

$$\left\| \nabla f(\theta^{k+1}) - \frac{1}{h^k}(\theta^{k+1} - \theta^{k+1/2}) \right\| \leq \epsilon$$

- ▶ analog of  $\|\nabla f(\theta^{k+1})\| \leq \epsilon$  for gradient method
- ▶ second term  $-\frac{1}{h^k}(\theta^{k+1} - \theta^{k+1/2})$  serves the purpose of a gradient for  $g$  (which need not be differentiable)

## Prox-gradient method summary

choose an initial  $\theta^1 \in \mathbf{R}^d$  and  $h^1 > 0$  (e.g.,  $\theta^1 = 0$ ,  $h^1 = 1$ )

for  $k = 1, 2, \dots, k^{\max}$

1. gradient step.  $\theta^{k+1/2} = \theta^k - h^k \nabla f(\theta^k)$
2. prox step.  $\theta^{\text{tent}} = \operatorname{argmin}_{\theta} \left( g(\theta) + \frac{1}{2h^k} \|\theta - \theta^{k+1/2}\|^2 \right)$
3. if  $F(\theta^{\text{tent}}) \leq F(\theta^k)$ ,
  - (a) set  $\theta^{k+1} = \theta^{\text{tent}}$ ,  $h^{k+1} = 1.2h^k$
  - (b) quit if  $\left\| \nabla f(\theta^{k+1}) - \frac{1}{h^k} (\theta^{k+1} - \theta^{k+1/2}) \right\| \leq \epsilon$
4. else set  $h^k := 0.5h^k$  and go to step 1



## Prox-gradient method convergence

- ▶ prox-gradient method always finds a stationary point
  - ▶ suitably defined for non-differentiable functions
  - ▶ assuming some technical conditions hold
- ▶ for *convex problems*
  - ▶ prox-gradient method is *non-heuristic*
  - ▶ for any starting point  $\theta^1$ ,  $F(\theta^k) \rightarrow F^*$  as  $k \rightarrow \infty$
- ▶ for *non-convex problems*
  - ▶ prox-gradient method is *heuristic*
  - ▶ we can (and often do) have  $F(\theta^k) \not\rightarrow F^*$

## Prox-gradient for regularized ERM

## Prox-gradient for sum squares regularizer

- ▶ let's apply prox-gradient method to  $F(\theta) = \mathcal{L}(\theta) + \lambda \|\theta\|_2^2$ 
  - ▶  $f(\theta) = \mathcal{L}(\theta)$
  - ▶  $g(\theta) = \lambda \|\theta\|_2^2 = \lambda \theta_1^2 + \dots + \lambda \theta_d^2$
- ▶ in prox step, we need to minimize  $\lambda \theta_i^2 + \frac{1}{2h^k}(\theta_i - \theta_i^{k+1/2})^2$  over  $\theta_i$
- ▶ solution is  $\theta_i = \frac{1}{1+2\lambda h^k} \theta_i^{k+1/2}$
- ▶ so prox step just shrinks the gradient step  $\theta^{k+1/2}$  by the factor  $\frac{1}{1+2\lambda h^k}$
- ▶ prox-gradient iteration:
  1. gradient step:  $\theta^{k+1/2} = \theta^k - h^k \nabla \mathcal{L}(\theta^k)$
  2. prox step:  $\theta^{k+1} = \frac{1}{1+2\lambda h^k} \theta^{k+1/2}$

## Prox-gradient for $\ell_1$ regularizer

- ▶ let's apply prox-gradient method to  $F(\theta) = \mathcal{L}(\theta) + \lambda \|\theta\|_1$ 
  - ▶  $f(\theta) = \mathcal{L}(\theta)$
  - ▶  $g(\theta) = \lambda \|\theta\|_1 = \lambda |\theta_1| + \dots + \lambda |\theta_d|$
- ▶ in prox step, we need to minimize  $\lambda |\theta_i| + \frac{1}{2h^k} (\theta_i - \theta_i^{k+1/2})^2$  over  $\theta_i$


- ▶ solution is

$$\theta_i^{k+1} = \begin{cases} \theta_i^{k+1/2} - \lambda h^k & \theta_i^{k+1/2} > \lambda h^k \\ 0 & |\theta_i^{k+1/2}| \leq \lambda h^k \\ \theta_i^{k+1/2} + \lambda h^k & \theta_i^{k+1/2} < -\lambda h^k \end{cases}$$

- ▶ called *soft threshold function*
- ▶ sometimes written as

$$\begin{aligned} \theta_i^{k+1} &= S_{\lambda h^k}(\theta_i^{k+1/2}) = \text{sign}(\theta_i^{k+1/2}) (|\theta_i^{k+1/2}| - \lambda h^k)_+ \\ &= (\theta_i^{k+1/2} - \lambda h^k)_+ - (-\theta_i^{k+1/2} - \lambda h^k)_+ \end{aligned}$$

## Soft threshold function



► prox-gradient iteration for regularized ERM with  $\ell_1$  regularization:

1. gradient step:  $\theta^{k+1/2} = \theta^k - h^k \nabla \mathcal{L}(\theta^k)$
2. prox step:  $\theta_i^{k+1} = S_{\lambda h^k}(\theta_i^{k+1/2})$  for  $i = 1, \dots, d$ .

- the soft threshold step shrinks all coefficients
- and sets the small ones to zero

## Prox-gradient step for nonnegative regularizer

- ▶ let's apply prox-gradient method to  $F(\theta) = \mathcal{L}(\theta) + r(\theta)$ , where  $r(\theta) = 0$  for  $\theta \geq 0$ ,  $\infty$  otherwise
  - ▶  $f(\theta) = \mathcal{L}(\theta)$
  - ▶  $g(\theta) = q(\theta_1) + \dots + q(\theta_d)$
- ▶ in prox step, we need to minimize  $q(\theta_i) + \frac{1}{2h^k}(\theta_i - \theta_i^{k+1/2})^2$  over  $\theta_i$
- ▶ solution is  $\theta_i = \left(\theta_i^{k+1/2}\right)_+$
- ▶ so prox step just replaces the gradient step  $\theta_i^{k+1/2}$  with its positive part
- ▶ prox gradient iteration:
  1. gradient step:  $\theta^{k+1/2} = \theta^k - h^k \nabla \mathcal{L}(\theta^k)$
  2. prox step:  $\theta^{k+1} = \left(\theta^{k+1/2}\right)_+$